

Irrigation Makes It Rain? Climate Spillovers and Adaptation in India

Mathilde Degois*

September 21, 2026

The most recent version is available [here](#).

*Preliminary version. Please, **do not share**.*

Abstract

Can economic development change the climate and reshape its own trajectory? I study how large-scale irrigation infrastructure reshapes weather patterns and agricultural development in India. Combining five decades of administrative, climatic, and agricultural data, I exploit wind direction as a source of exogenous exposure to distant irrigation schemes. A 10 p.p. increase in irrigation, about half of India's Green Revolution expansion, raises dry-season rainfall by 5 mm per month and lowers temperatures by 0.3°C. Effects reverse during the monsoon, with declining rainfall and seven additional days of crop-damaging heat. Crop production rises by 5% in the dry season without yield gains, whereas wet-season crop yield increases by 1.3%. This puzzling result points to adaptation rather than solely a direct climate effect on agriculture. Farmers expand private irrigation downwind, enabling double cropping on more marginal land. In the long run, these spillovers sustain smallholders while enabling larger farms to expand, polarizing the agricultural sector. Overall, these climatic spillovers account for 7% of crop production growth during the period, worth an extra \$1 billion a year.

Keywords: Irrigation, Climate Spillovers, Adaptation, Climate Change, India

JEL codes: O13, 053, Q15, Q54, Q56, Q58.

*Email: mathilde.degois@ens-paris-saclay.fr. ENS Paris-Saclay, CEPS & AgroParisTech, PSAE.

I thank Tamma Carleton, Kelsey Jack, William Boos, Sandy Sum, Anna Papp, Ram Fishman, Simon Greenhill, Marco Gonzalez-Navarro, Ted Miguel, Michael Anderson, Ellen Bruno, Eugenie Dugoua, Florian Grosset-Touba, Pierre Mérel, Hugo Molina, Yannick Guyonvarch, Julien Wolfersberger, and Thomas Vendryes for their insights and helpful comments. I thank Tamma Carleton's team for excellent feedback. I thank seminar and conference participants from Development and Environment PhD Workshop at OsloMet, PSAE Internal Seminar, IPWSD Workshop, Climate Economic Lunch ARE at UC Berkeley, Agricultural Econ. PhD Workshop ARE at UC Davis, Big Sky Workshop, AMIE Workshop, WCERE, FAERE, Land-Use Policies & Climate Change Adaptation Workshop and BioEcon. Any remaining errors are my own.

1 Introduction

Irrigation is a cornerstone of agricultural development, covering one-fifth of the world’s cropland (FAO, 2023). It raises productivity, reduces poverty (Dillon and Fishman, 2019; Bassini et al., 2019), and helps farmers adapt to climate risk (Taylor, 2025). Yet irrigation also generates environmental externalities. Intensive use drives rapid groundwater depletion (Wada et al., 2010), which in turn constrains development (Sekhri, 2014; Blakeslee et al., 2020). Beyond resource depletion, the climate science literature documents that irrigation alters the climate, affecting rainfall and temperature beyond irrigated areas (Puma and Cook, 2010; Alter et al., 2015; Nocco et al., 2019; Greve et al., 2025). By increasing evapotranspiration, irrigation modifies land–atmosphere interactions, affecting cloud formation and precipitation downwind. What remains largely unexplored is the empirical assessment of these climatic spillovers, and, more importantly, their consequences on economic development.

This paper addresses this gap by answering two linked questions: (1) to what extent does irrigation infrastructure alter the climate? and (2) through this channel, what are the short- and long-run consequences for development?

I study these questions in the context of India, the country with the largest irrigated area in the world, over the period 1960–2020. Over the past six decades, India has built an extensive network of dams and canals that now irrigates nearly one-third of its cropped area. These investments were rolled out in a staggered fashion across space, generating substantial spatial and temporal variation. During the same period, the country has experienced a delayed structural transformation, with a large share of the population remaining in agriculture and a slow transition toward larger and more efficient farm structures.

To examine how irrigation affects both climate and agricultural development, I combine historical administrative records on irrigation infrastructure with high-resolution climate data and administrative agricultural data. Causal identification exploits exogenous variation in exposure to distant upwind irrigation projects generated by wind flows, while absorbing local, seasonal, and state-year shocks. I trace the resulting responses at short, medium, and long horizons using

annual panels, five- and ten-year averages, and long differences. The results are not driven by potential wind-borne confounders, such as air pollution, nor by the Green Revolution, and are robust to placebo specifications that exploit reversed wind direction.

I first demonstrate that large-scale surface irrigation infrastructure alters the climate. A 10 percentage point increase in upwind irrigation exposure, about half of India's Green Revolution expansion, raises annual precipitation and lowers high temperatures. Yet these aggregate effects mask sharp seasonal asymmetries. During the dry Rabi season, exposure increases rainfall by 6–7 mm per month and cools temperatures by about 0.5°C. During the wet Kharif season, the effect reverses: rainfall falls by a similar magnitude, monsoon onset is delayed, and exposure to harmful heat above 29°C rises by six days in June. These climatic spillovers persist and strengthen in time.

These seasonal changes have theoretically ambiguous effects on agricultural production and adaptation. While higher annual rainfall would typically raise yields and reduce irrigation incentives, its seasonal distribution matters. Lower monsoon rainfall increases the value of irrigation as both a substitute for rainfall and insurance against growing-season risk, while higher dry-season rainfall raises its return by making second-season cultivation more profitable. Accounting for seasonal redistribution therefore highlights a channel through which rainfall and irrigation may act as complements rather than substitutes.

Building on these insights, and having established that irrigation modifies downwind climatic conditions, I next examine its implications for agricultural development. A 10 percentage point increase in upwind irrigation exposure raises cereal production by 5% for wheat, the major dry-season crop in India, yet with no improvement in yields. Rice, the major wet season crop, observe an slight 1.3% increase of yield, with no significant impact on production or area cultivated. This symmetry is striking because the two crops face opposite seasonal shocks: dry-season conditions improve for wheat, while monsoon conditions deteriorate for rice.

Therefore, equal production gains from opposing shocks point to adaptation, rather than direct productivity gains, as the key margin. Consistent with theoretical predictions, farmers

respond by increasing irrigation intensity by 1.3 percentage points, about 5% of India's aggregate increase between 1960 and 2010. This adaptation operates differently across seasons. In the dry season, improved climatic conditions, together with greater irrigation access, make second-season cultivation more viable. Cropping intensity increases by 1.2 percentage points, which represents roughly 4% of India's total gain over the past half century. Since wheat yields do not increase, this response appears to operate mainly through the expansion of cultivation onto more marginal land.¹ During the monsoon, irrigation plays a different role. As rainfall declines, temperature increases and the onset of the season becomes more uncertain, irrigation buffers farmers against adverse climatic conditions: rice production do not fall as farmers are able to sustain land productivity. In contrast to evidence from developed agricultural settings such as the United States (Grosset-Touba et al., 2024), I find no to negative evidence that farmers respond by investing in other adaptation margins, such as mechanization, fertilizer use, or major changes in crop composition.

Investigating the mechanism behind irrigation adoption, I find that the response is driven both by substitutability between irrigation and rainfall during the wet season and complementarity enabling dry-season production expansion. These results highlight the importance of accounting for seasonal climatic variation in farmers' adaptation decisions, which can have important implications for long-run agricultural investments.

Finally, irrigation-induced weather variation reshapes India's farm structure. On the one hand, I observe the sustainability of the survival of small farms, slowing transitions toward medium-scale holdings. On the other hand, large farms expand their operated area without a corresponding increase in their number. This results in a polarization within agriculture: promoting consolidation at the top and stagnation at the bottom, delaying agricultural structural transformation.

A counterfactual exercise estimates that between 1967 and 2020, upwind irrigation's climatic

¹In India, increases in cropping intensity have been an important margin for raising farm incomes and reducing rural poverty, and have contributed substantially to food production gains over the past 50 years. See <http://vsliir.iima.ac.in:8080/xmlui/handle/11718/17078>.

spillovers account for 7 percent of India’s wheat production growth over the Green Revolution period, worth an extra 1 billion USD a year. Ongoing construction projects are projected to irrigate an additional 32.3 million hectares, implying 48 km³ of additional downwind dry-season rainfall (8.4 percent of India’s observed dry-season rainfall) and 3.6 million tonnes of wheat per year (4.6 percent of current annual production), with 30% of those spillovers falling on rainfed land.

Related literature. Focusing on irrigation, this paper improves our understanding of the feedbacks between development and climate (Jack and Ryan, 2026). It contributes to three strands of the literature.

First, this paper provides new causal evidence on the impacts of irrigation infrastructure on development by documenting a previously overlooked externality, namely climate spillovers. A large body of work has studied irrigation’s economic benefits since the seminal work of Duflo and Pande (2007), with subsequent studies showing substantial gains in farm income, productivity, and resilience from surface irrigation (Hansen et al., 2011; Strobl and Strobl, 2011; Blanc and Strobl, 2014; Olmstead and Sigman, 2015; Dillon and Fishman, 2019; Zaveri et al., 2020; Jones et al., 2022).² More recent work emphasizes local productivity gains within irrigated areas and their implications for structural transformation (Asher et al., 2022; Blakeslee et al., 2023).³ This paper adds to this literature by showing that irrigation infrastructure reshapes climatic conditions through the spatial redistribution of rainfall and temperature downwind, and by quantifying how this climatic shift affects agricultural production, farmers’ investment and adaptation strategies, and long-run agricultural structure over five decades.

The second contribution bridges climatological and economic evidence by showing that climate is partly endogenous to human activity, with important implications for economic development. Natural science evidence shows that land-use changes – deforestation, irrigation, crop

²Other studies document similar effects for groundwater irrigation (Fishman et al., 2013; Sekhri, 2014; Hornbeck and Keskin, 2014, 2015; Ryan and Sudarshan, 2020; Blakeslee et al., 2020; Butler and Boudot Reddy, 2023).

³More broadly, this paper also relates to the literature on the spatial allocation of infrastructure and place-based policies (Donaldson and Hornbeck, 2016; Donaldson, 2018; Asher and Novosad, 2020; Bustos et al., 2016, 2020).

choice, urbanization – alter land–atmosphere interactions, with documented effects on temperature and precipitation (Puma and Cook, 2010; Alter et al., 2015; Greve et al., 2025; Butt et al., 2011; Leite-Filho et al., 2021; Loarie et al., 2011; Manoli et al., 2019).⁴ A nascent growing literature in economics has begun to examine these effects. Focusing on forest cover change, it documents the downwind consequences of deforestation for hydropower in the Amazon (Araujo, 2023), conflict in West Africa (Degois et al., 2025), and of U.S. afforestation programs for agriculture (Li, 2021; Grosset-Touba et al., 2024). Closely related to this paper, Braun and Schlenker (2023) document that irrigation in the United States lowers temperatures, increasing agricultural productivity and reducing mortality. This paper shows that the climatic effects of irrigation depend on the underlying climate regime. In contrast to existing evidence from temperate regions, I find that irrigation reshapes the seasonal distribution of rainfall in monsoon systems, which sustain nearly two-thirds of the world’s population. Climate models have predicted such responses (Saeed et al., 2009; Guimberteau et al., 2012a; Chou et al., 2018; Devanand et al., 2019); this paper provides, to my knowledge, the first empirical evidence. I then trace the economic consequences of this climate feedback over the short, medium, and long run.

Third, this paper contributes to the literature on climate change, adaptation, and development. I document an empirical feedback loop between adaptation and climate through the case of irrigation, and introduce a new source of quasi-experimental variation to study the effects of persistent climatic change. Existing work typically relies either on short-run weather shocks (Schlenker and Roberts, 2009; Lobell et al., 2011; Dell et al., 2012, 2014; Carleton and Hsiang, 2016) or on long-run differences in climate trends across locations (Henderson et al., 2017; Kala, 2017; Taraz, 2017, 2018; Emerick, 2018; Liu et al., 2023) to study the impact of climate change on the economy. Both approaches treat climate variation as external to local adaptation decisions (Dell et al., 2014). Yet farmers are known to adapt to weather shocks and monsoon shifts through irrigation investment (Hrozencik et al., 2025; Taraz, 2017; Kala, 2017). If irrigation also affects rainfall and temperature, observed climate variation may partly embed past adaptation responses. This

⁴Go to [te Wierik et al. \(2021\)](#) for a review of the literature.

paper provides an empirical evidence of this feedback: irrigation is both a driver of persistent climatic change and an adaptation to it.⁵ I address this challenge by exploiting irrigation-induced climatic spillovers documented in this paper, which generate climate variation that is persistent yet orthogonal to local adaptation decisions. Finally, the paper shows that understanding adaptation requires looking beyond annual climate averages to the seasonal distribution of rainfall and temperature.

The remainder of the paper is organized as follows. Section 2 presents the climatic and agricultural context. Section 3 presents a simple model of irrigation adoption under seasonal rainfall reallocation, to build intuition for the results that follow. Section 4 describes the data sources and variables. Section 5 outlines the empirical strategy. Section 6 presents the main results, Section 7 discusses robustness, and Section 9 concludes.

2 Climatic and Agricultural Context

This section describes the climatic and agricultural context of the study, characterized by the Indian monsoon system, the agricultural calendar, and irrigation use.

The Indian Monsoon India’s climate is dominated by the Southwest Monsoon, delivering 70 percent of India’s annual rainfall over the months of June to September – on average around 870 mm, based on the 1971–2020 long-period average (IMD, 2026). The rest of the year is particularly dry, especially March to May, with an average of around 130 mm of rain over these three months. Appendix Figure A7 plots this seasonal cycle for the districts of the estimation sample: the median district receives 288 mm in July against 10 mm in January. The two seasons are not scaled versions of one another in space either. Appendix Figure A8 shows that monsoon rainfall is heaviest on the Western Ghats and in the north-east, while dry-season rainfall falls in the north-west and

⁵This has methodological implications: if adaptation has moderated climatic extremes, estimates based on observed long-run climate variation may understate true damages, generating attenuation bias. This concern relates to the interpretation of long-run climate estimates as evidence of adaptation (Burke and Emerick, 2016; Liu et al., 2023) and aligns with recent macro-climate models that endogenize adaptation in spatial equilibrium dynamics (Krusell and Smith, 2022; Cruz and Rossi-Hansberg, 2024; Fillon, 2025).

along the south-eastern coast, which receives the retreating monsoon. Temperature follows the same calendar, with a median district-month of 21.4°C in Rabi against 27.2°C during the monsoon (Appendix Figure A9). Over the past two decades, monsoon onset and progression have grown increasingly erratic, with the underlying causes still debated in the natural science literature (Kapa et al., 2025).

The Agricultural Calendar: Kharif and Rabi India’s agricultural calendar follows this seasonal rhythm and is organized into two main cropping seasons. The *Kharif* season runs from June to September, coinciding with the monsoon: rain-fed crops such as rice, cotton, and pulses are sown at the onset of the rains and harvested in autumn. The *Rabi* season runs from roughly November to March, during the dry winter months, and is dominated by wheat. Because little rain falls during this period, Rabi cultivation depends almost entirely on irrigation. Access to irrigation therefore determines whether farmers can grow a second crop each year, a key margin for raising agricultural income (Duflo and Pande, 2007). Irrigation is also used during the Kharif season itself, despite ongoing rainfall: across the Indo-Gangetic Plain, farmers commonly rely on supplementary, “life-saving” irrigation to bridge dry spells within an increasingly uneven monsoon (Urfels et al., 2020).

Sources and Cost of Irrigation Irrigation in India is delivered through two main channels: canal networks fed by dams and reservoirs, and groundwater extracted through privately owned tubewells. Canals were the principal source for most of the period studied here, before tubewells overtook them as the dominant source from the 1990s onward; groundwater now accounts for the majority of irrigated area nationally (Singh et al., 2022). The two technologies differ in cost structure: canal infrastructure requires large, lumpy public investment, while a tubewell requires a smaller, sunk drilling cost, after which pumping is cheap, especially given subsidized agricultural electricity. Tubewell use also varies by season, with utilization roughly 7.5 times higher during Kharif than during Rabi.

3 Irrigation Adoption under Seasonal Rainfall Reallocation

If irrigation is not only a productive input but also reshapes the seasonal distribution of rainfall downstream, an important question arises: how do such climatic spillovers affect other farmers' incentives to invest in irrigation? Potentially creating feedback effects across space and time.

In standard theoretical models, irrigation and rainfall are technological substitutes. When rainfall is abundant, the marginal return to irrigation declines because water inputs exhibit diminishing returns (Feder et al., 1985). In such frameworks rainfall enters production as a single aggregate input, and irrigation augments total water availability. However, in many agricultural systems irrigation does more than increase productivity within a growing season. In environments with a distinct monsoon and dry season, irrigation may enable cultivation during the dry season and increase cropping intensity (Jones et al., 2022; Hussain and Hanjra, 2004; Foster and Rosenzweig, 1995; Duflo and Pande, 2007). In these contexts irrigation affects production along two margins: it increases water availability within season and expands the feasible production calendar by activating an additional cropping season. This distinction implies that the relationship between rainfall and irrigation adoption depends not only on total annual rainfall but also on its seasonal distribution.

This section develops a simple model of irrigation adoption under seasonal rainfall variation. The model highlights how changes in the timing of rainfall can generate different adoption responses in contrast to aggregated annual totals.

3.1 Baseline model: Single-Season Agriculture

I begin with a standard one-season model to establish the benchmark prediction. A farmer cultivates one crop during the monsoon season. Land and other inputs are fixed. The farmer is risk neutral and chooses whether to invest in irrigation. Installing irrigation requires a fixed cost $F > 0$ and provides an additional water input $\alpha > 0$.

Let rainfall be denoted by r . Effective water is given by total water availability $w = r + \alpha I$,

where $I \in \{0, 1\}$ denotes irrigation adoption. Output per hectare is $f(w)$, where $f'(w) > 0$ and $f''(w) < 0$, implying diminishing returns to water.

Profits are given by $\Pi(I; r) = f(r + \alpha I) - FI$. The gain from adopting irrigation is therefore

$$\Delta(r) = f(r + \alpha) - f(r) - F. \quad (1)$$

Because production is concave in water, the marginal product of water declines as rainfall increases.

Proposition 1 (Substitution). *Higher rainfall reduces the profitability of irrigation:*

$$\frac{\partial \Delta}{\partial r} < 0.$$

Thus in the standard single-season environment rainfall and irrigation are technological substitutes, and rainfall shocks reduce incentives to adopt irrigation.

3.2 Two-Season Agriculture

Now consider an agricultural year composed of two seasons: the monsoon season (Kharif), indexed K , and the dry season (Rabi), indexed R . Seasonal rainfall is denoted by (r_K, r_R) , with total annual rainfall $r = r_K + r_R$.

A key feature of the environment is that dry-season cultivation requires irrigation.⁶ This assumption captures environments where residual soil moisture is insufficient to sustain Rabi production. Irrigation therefore does not merely increase productivity; it activates production in an additional season.

Output without irrigation is $\Pi^0(r_K) = f_K(r_K)$. When irrigation is adopted, both seasons are cultivated and total output is

$$\Pi^1(r_K, r_R) = f_K(r_K + \alpha_K) + f_R(r_R + \alpha_R) - F,$$

where $f_s(\cdot)$ is increasing and concave in effective water and α_s denotes irrigation water in season s . Irrigation is adopted whenever the gain $\Delta(r_K, r_R) = \Pi^1 - \Pi^0$ is non-negative. Rewriting the

⁶More generally, irrigation may increase the share of land that can be cultivated during the dry season.

gain yields

$$\Delta(r_K, r_R) = \underbrace{[f_K(r_K + \alpha_K) - f_K(r_K)]}_{\text{Substitution}} + \underbrace{f_R(r_R + \alpha_R)}_{\text{Activation}} - F. \quad (2)$$

Two distinct channels determine irrigation profitability. The first term captures the *substitution effect*: irrigation increases water availability during the monsoon season when rainfall is scarce. The second term captures the *activation effect*: irrigation enables cultivation during the dry season. Seasonal rainfall affects irrigation adoption through these two margins.

Following Proposition 1, higher rainfall during the monsoon season reduces irrigation profitability ($\frac{\partial \Delta}{\partial r_K} < 0$) as it lowers the marginal value of irrigation. What we are interested in is the second activation term that leads to a positive effect of higher Rabi rainfall on irrigation adoption.

Proposition 2 (Activation). *Higher rainfall during the dry season increases irrigation profitability:*

$$\frac{\partial \Delta}{\partial r_R} > 0.$$

Higher rainfall during the dry season (r_R) raises the productivity of Rabi cultivation, which is activated only when irrigation is installed.

The overall effect of rainfall shocks therefore depends on how rainfall is distributed across seasons. Consider a change in rainfall allocation such that total rainfall changes by $\Delta r = \Delta r_K + \Delta r_R$. Irrigation adoption increases whenever

$$f'_R(r_R + \alpha_R)\Delta r_R > [f'_K(r_K) - f'_K(r_K + \alpha_K)]|\Delta r_K|. \quad (3)$$

Climatic shocks that reallocate rainfall across seasons may therefore either increase or decrease irrigation incentives depending on their direction, the degree of substitution between rainfall and irrigation during the monsoon season, and the profitability of dry-season production.

The comparative statics above can be summarized by an adoption frontier that separates rainfall states in which irrigation is profitable from those in which it is not. Figure A12 in Appendix B illustrates the irrigation adoption frontier implied by the model. Irrigation is adopted whenever the pair (r_K, r_R) lies above the frontier. Panels (b)–(d) illustrate how the adoption frontier shifts. Higher Rabi productivity lowers the threshold because irrigation becomes more valuable

through the activation of dry-season production. A higher Kharif scarcity value increases the marginal benefit of irrigation when monsoon rainfall is low. Finally, higher fixed costs shift the entire frontier upward, reducing irrigation adoption in all rainfall states.

In contrast to the single-season benchmark where higher total rainfall would unambiguously reduce irrigation incentives, the two-season environment underlines that the effect depends on how rainfall is distributed across seasons.

4 Data

The analysis combines spatial and temporal variation in irrigation coverage across India with remote sensing measures of climate and district-level agricultural panel data. I first estimate the effects of irrigation on downwind climate using district-month observations from 1966 to 2020, allowing me to account for the seasonal dynamics of atmospheric moisture transport. Because the agricultural data are available only at the annual level, I then aggregate the irrigation exposure measure accordingly to examine how irrigation-induced precipitation shapes downstream agricultural productivity and farmers' adaptation. Appendix Table A7 reports descriptive statistics for the estimation sample.

Climate Data. To characterize atmospheric dynamics, this paper uses the ERA5 reanalysis dataset developed by the European Centre for Medium-Range Weather Forecasts (ECMWF).⁷ ERA5 provides pressure fields together with the horizontal wind components (u and v), which jointly describe wind speed and direction. These data are available globally at a spatial resolution of 0.25 degree (approximately 24 km at the equator) at monthly frequency. The analysis focuses on the 1960–2020 period, during which wind trajectories are reconstructed by tracing air mass movements backward in time in order to connect pixels, following the approach of Araujo (2023) for each pixel-month. Rainfall data primarily come from gridded precipitation data produced by the India Meteorological Department (IMD), which reports daily rainfall at a resolution of 0.25

⁷Downloaded from ERA5 on the 17/06/2025.

degree for the period 1960–2024 (Pai et al., 2014).⁸ This dataset is widely used in the literature (Burgess et al., 2017) and considered the most reliable source for India.⁹ Temperature data come from ERA5-Land, available at a finer resolution of 0.1 degree (about 9 km at the equator).

Irrigation Data. Information on irrigation infrastructure is obtained from the India Water Resources Information System (WRIS), which provides GIS data on canals, command areas, aquifers, and rivers.¹⁰ An irrigation project typically consists of a reservoir (or a river diversion) that supplies a network of canals distributing water to the surrounding area. The command area of a project is defined as the land that can be reached by these canals and is therefore eligible for irrigation. The size of the command area is determined by the storage capacity of water (often a dammed reservoir, though sometimes a direct river diversion) and the terrain’s topography. In addition to spatial coverage, I retrieve information on the start and end years of canal construction as well as project status (ongoing vs. completed) using the Central Water Commission’s Management Information System of Water Resources Projects and the India WRIS Wiki. As documented by Asher et al. (2022), reported canal start dates sometimes refer to rehabilitation rather than initial construction. To address this issue, open-access construction dates compiled by Asher et al. (2022) are used and supplemented with manual validation based on government documents, news articles, and academic reports.¹¹ Figure 1 maps the spatial and temporal coverage of canal irrigation projects (i.e., command areas) since 1960.¹² The irrigation variable is constructed as the proportion of each pixel’s area irrigated by canals in a given year. This measure, which ranges from 0 to 1, captures the cumulative expansion of irrigated command areas over time at the pixel

⁸Dataset available at <https://imdpune.gov.in/cmpg/Griddata/Rainfall25Bin.html>.

⁹As a robustness check, rainfall estimates from ERA5-Land are also used. Results are consistent across both data sources.

¹⁰Canal shapefiles and identifiers were originally available at <https://nwdp.nwic.in/dataset/canal>, accessed on 28/05/2025. The dataset was last updated on May 5, 2025 (10:46 AM, UTC+02:00), but the link is no longer accessible.

¹¹I acknowledge the use of ChatGPT’s deep search to support this validation, with all findings cross-checked against primary sources.

¹²Table A8 shows that canals with an unknown completion year represent 3.5% of all completed canals in the sample, and that approximately 18% of India’s total area is currently irrigated by canals. An additional 8% of the country’s surface area is under ongoing canal construction projects. I perform a forecasting analysis on future canal construction in the following.

level.




Figure 1: Command Area Construction Years across India

Note: Figure 1 illustrates the canal-irrigated command areas included in the sample, displayed by their year of completion for all available commands.

Agricultural Data. Agricultural outcomes are measured using the District-Level Data (DLD) from ICRISAT, which provides a comprehensive yearly panel of agricultural statistics for Indian districts from 1966 to 2011. District boundaries correspond to those of India in 1966, including 311 districts at the time.¹³ Figure A11 in the Appendix maps upwind irrigation exposure on these 1966 boundaries, in level and in within-district variation. The dataset covers a wide range of agricultural indicators, including cropped area, production and yields (calculated as production over area) disaggregated by major crop types. In addition, it includes information on agricultural practices such as irrigation by type, fertilizer use, high-yielding variety (HYV) crop adoption, double cropping, and fallow land. Data on farm size distribution are available only at five-year intervals, corresponding to the Agricultural Census of India for census years from 1970 to 2005. These data report the number and total area of farms disaggregated into five size categories: *marginal* (< 1 ha), *small* (1–2 ha), *semi-medium* (2–4 ha), *medium* (4–10 ha), and *large* (> 10 ha).

Additional Controls. The robustness analysis in Section 7 draws on a broad set of geographic, agricultural, and environmental covariates. Geographic controls include elevation and its within-district dispersion, slope, terrain ruggedness, distance to the coast, latitude, longitude, district area, and aquifer composition. I construct the topographic measures from WorldClim’s 30 arc-second SRTM-derived elevation data (Fick and Hijmans, 2017), using the terrain ruggedness index of Riley et al. (1999), and derive aquifer shares from the Central Ground Water Board’s principal aquifer map. Agricultural controls include baseline forest cover and irrigation, population, and the area planted with high-yielding varieties. Finally, I measure air pollution using the monthly satellite-derived surface PM_{2.5} product of van Donkelaar et al. (2021), available from 1998 at a 0.01° resolution and aggregated to the 0.25° analysis grid.

¹³ICRISAT consolidated data on Indian districts from multiple sources, see [ICRISAT Data](#) and [Shapefile](#).

5 Empirical Strategy

The empirical strategy exploits variation in prevailing wind trajectories and the timing of irrigation expansion to construct quasi-experimental exposure to upwind irrigation. The underlying idea is that air masses passing over irrigated land acquire additional moisture, thereby affecting rainfall and temperature downwind and, ultimately, agricultural outcomes. Identification relies on the exogeneity of wind paths, rich spatial and temporal fixed effects, and the substantial distance between focal cells and their upwind counterparts, which limits the scope for reverse causality. I first describe the construction of the exposure measure in Section 5.1. I then present the empirical framework used to estimate its relationship with rainfall and temperature in Section 5.2, before extending the analysis to agricultural outcomes in Section 5.3. Finally, Section 5.4 outlines the empirical approach used to study the medium- and long-run effects.

5.1 Methodology


The key explanatory variable, *Upwind Irrigation*, is constructed by combining wind-trajectory with the spatial irrigation dataset described in Section 4. The following wind connectivity matrix is defined at a 0.25° spatial resolution (approximately 24 km at the equator). For each pixel of interest i , I (i) trace its monthly backward wind path at the 800 hPa level,¹⁴ and (ii) retain only the segments of this trajectory that lie between 50 and 150km from i , following Araujo (2023) and Degois et al. (2025). Pixels within 50 km are excluded to minimize potential reverse causality (e.g., irrigation decisions in neighboring areas could be correlated with local climate conditions), while segments beyond 150km are discarded because the influence of irrigation on moisture transport is expected to dissipate at larger distances (Spracklen et al., 2012; Degois et al., 2025).¹⁵ For each pixel i , year t , and month m , *Upwind Irrigation* is defined as the average share of irrigated land across cells lying along i 's backward wind trajectory within the 50–150 km window. By construction,

¹⁴This choice ensures that wind paths are determined at an altitude largely unaffected by local surface conditions such as vegetation or irrigation, and can therefore be considered exogenous to pixel-level agricultural activities. In addition, this is the atmospheric level at which most clouds form, strengthening our ability to capture the effects of rainfall rather than other wind-borne confounders.

¹⁵Robustness checks using alternative exclusion buffers and outer radii do not change the results.

this variable takes values in $[0, 1]$ and varies over time due to year-to-year changes in irrigated area and monthly fluctuations in wind direction.

Figure 2: Upwind Irrigation



Notes: The **upwind neighbors** (in red) are linked to the **pixel of interest i** (in bold black) through its backward wind path, represented in blue. To mitigate concerns about reverse causality, pixels located within a 50 km buffer around the centroid of the pixel of interest (in green) are excluded, as are pixels located more than 150km away. The explanatory variable *Upwind Irrigation* is constructed by identifying all pixels lying along the backward wind path and located between 50 and 150km from each pixel of interest (highlighted in red). *Upwind Irrigation* is then defined as the average proportion of irrigated area across these upwind neighbors. The figure is not drawn to scale and is intended for illustrative purposes only.

Figure 2 illustrates this construction. The pixel of interest i (in black) is connected to its upwind neighbors (in red) through its backward wind path (in blue). Neighbors along the wind path falling within the 50–150km window are retained, while those too close (in green) or too distant are excluded. The resulting measure thus captures the intensity of irrigation in the relevant upwind zone. Figure 3 illustrates the independent variable *Upwind Irrigation* across India, averaged over the period 1966–2020. To implement placebo tests, I reverse the connection matrix and repeat the same procedure, defining exposure to downwind rather than upwind irrigation.

District-Level Exposure. Because agricultural outcomes are observed at the district level, I construct a district-level measure of *Upwind Irrigation* by spatially aggregating pixel-level ex-

posure. For each district, year, and month, I average exposure across all pixels intersecting the district, using the area of overlap between each pixel and the district as spatial weights. Formally,

$$E_{d,t,m}^{\uparrow} = \frac{\sum_{i \in d} e_{i,t,m}^{\uparrow} a_{i \cap d}}{\sum_{i \in d} a_{i \cap d}},$$

where $e_{i,t,m}^{\uparrow}$ is the monthly grid-cell-level *Upwind Irrigation* exposure of pixel i in year t and month m , and $a_{i \cap d}$ represents the area of overlap between pixel i and district d . This procedure yields a continuous monthly district-level exposure measure bounded between 0 and 1, which is comparable across districts and independent of district size. To avoid mechanical endogeneity, irrigation located within the district itself is excluded from the calculation, such that exposure captures only irrigation occurring in upwind districts. The resulting district-level measure is plotted in Figure A11 for each of the 311 districts defined under the 1966 administrative boundaries included in the sample.

5.2 Irrigation and Climate

I first examine whether irrigation occurring upwind affects downwind climatic conditions, focusing on precipitation and temperature. This relationship constitutes the first stage of the empirical design, establishing whether irrigation expansion influences atmospheric moisture transport downwind. Given the strong seasonality of both irrigation and the climate in India, the atmospheric impacts of irrigation are expected to vary over the agricultural calendar. To capture this seasonal heterogeneity, I estimate the following district-month specification:

$$y_{d,t,m} = \beta_{Jul} E_{d,t,m}^{\uparrow} + \sum_{m \neq Jul} \beta_m \mathbf{1}\{M_{d,t,m} = m\} \cdot E_{d,t,m}^{\uparrow} + \mu_d + \delta_m + \lambda_{s,t} + \varepsilon_{d,t,m} \quad (4)$$

where $y_{d,t,m}$ is the outcome of interest—total precipitation (in mm) or mean temperature (in °C)—for district d , year t , and month m . The key explanatory variable, *Upwind Irrigation* $E_{d,t,m}^{\uparrow}$, measures the share of irrigated land located upwind of district d at time (t, m) . The specification introduces $\mathbf{1}\{M_{d,t,m} = m\}$, an indicator variable for each month m , with July as the reference




Figure 3: Upwind Irrigation Exposure Across India

Notes: The figure shows the spatial distribution of the key explanatory variable, *Upwind Irrigation*, at a 0.1 degree resolution. It captures the average share of irrigated land across upwind pixels located along the seven-day backward wind trajectory, restricted to a 50–150km upstream range of each pixel i . Command areas are outlined in white. Values are averaged over the full study period, spanning 1966–2020.

category. District fixed effects, μ_d , absorb time-invariant location characteristics, month fixed effects δ_m account for seasonality, and state-by-year fixed effects $\lambda_{s,t}$ capture local shocks and regional climate trends. Standard errors are clustered at the district level to account for spatial correlation. β therefore captures the total effect of upwind irrigation exposure on climate variables in July. For any other month m , the total effect is given by $\beta + \gamma_m$. To ease inter-

pretation, I scale the estimates to a 10 percentage point increase in upwind irrigation exposure, roughly half of India’s irrigation expansion during the Green Revolution. The resulting month-specific coefficients provide a dynamic perspective on how both the direction and magnitude of irrigation-induced climate effects evolve throughout the year.

5.3 Climatic Spillovers, Agriculture and Adaptation

Having established that upwind irrigation can alter weather patterns downwind, I next examine whether such irrigation-induced climatic changes translate into measurable changes in agricultural performance.

Annual Aggregation. Because agricultural outcomes are observed annually, I aggregate monthly district-level exposure to the yearly level. I compute a time-weighted average using the number of days in each month as weights, which accounts for unequal month lengths:

$$E_{d,t}^{\uparrow} = \frac{\sum_{m=1}^{12} E_{d,t,m}^{\uparrow} w_{t,m}}{\sum_{m=1}^{12} w_{t,m}},$$

where $w_{t,m}$ denotes the number of days in month m of year t . In robustness checks, I apply the same aggregation procedure separately to wet- and dry-season months to construct season-specific annual exposure measures.

Effects on Agricultural Production. I then estimate the effect of upwind irrigation exposure on agricultural outcomes using the following district-year panel specification:

$$Y_{d,s,t} = \beta E_{d,t}^{\uparrow} + \gamma \mathbf{X}_{d,t} + \mu_d + \lambda_{s,t} + \varepsilon_{d,s,t}, \quad (5)$$

where $Y_{d,s,t}$ denotes agricultural outcomes (e.g., total production, cultivated area, yields, etc.) in district d and year t . The variable $E_{d,t}^{\uparrow}$ is the yearly district-level exposure to upwind irrigation. The vector $\mathbf{X}_{d,t}$ includes log population and log high-yielding varieties, which proxy for time-varying demographic pressure and contemporaneous technological upgrading. District

fixed effects μ_d absorb time-invariant district characteristics, including soil quality, topography, and baseline agro-climatic conditions, while state-by-year fixed effects $\lambda_{s,t}$ capture policy, market, and climatic shocks common to districts within the same state and year. Standard errors are clustered at the district level.

5.4 Medium- and Long-Run Effects

Finally, I examine whether irrigation spillovers persist over longer horizons. This exercise is motivated by two considerations. First, I test whether irrigation generates persistent changes in downwind weather conditions, rather than only transitory year-to-year fluctuations. Second, agricultural responses may take time to materialize, as farmers gradually adjust input use, crop choices, irrigation adoption and production practices. I therefore complement the annual panel estimates with moving-average and long-difference specifications.

I first estimate medium-run specifications using five- and ten-year moving averages. For each district d and anchor year τ , I average outcomes, upwind irrigation exposure, and controls over windows centered around τ . Let $L \in \{5, 10\}$ denote the window length. For any district-year variable $Z_{d,t}$ entering the specification, define

$$\bar{Z}_{d,\tau}^{(L)} = \frac{1}{L} \sum_{k \in \mathcal{W}_L(\tau)} Z_{d,k},$$

where $\mathcal{W}_5(\tau) = \{\tau - 2, \dots, \tau + 2\}$ and $\mathcal{W}_{10}(\tau) = \{\tau - 4, \dots, \tau + 5\}$. I estimate the five-year specification for census anchor years $\tau \in \{1970, \dots, 2000\}$,¹⁶ and the ten-year specification for $\tau \in \{1970, 1980, 1990, 2000\}$. I then estimate

$$\bar{y}_{d,\tau}^{(L)} = \beta_1 \bar{E}_{d,\tau}^{\uparrow,(L)} + \gamma \bar{\mathbf{X}}_{d,\tau}^{(L)} + \mu_d + \lambda_{s,\tau} + \varepsilon_{d,\tau}. \quad (6)$$

The specification includes district fixed effects μ_d and state-by-anchor-year fixed effects $\lambda_{s,\tau}$. Identification comes from changes in medium-run upwind irrigation exposure within districts,

¹⁶These anchor years correspond to the availability of India's agricultural census, which allows me to estimate the same specification for farm-structure outcomes.

relative to other districts in the same state and anchor year.

I further estimate a long-difference specification comparing the beginning and end of the sample. I define changes as the difference between five-year averages centered on 2000 and 1970:

$\Delta Z_d = \bar{Z}_{d,2000}^{(2)} - \bar{Z}_{d,1970}^{(2)}$. I then estimate

$$\Delta y_d = \beta_1 \Delta E_d^\uparrow + \gamma \Delta \mathbf{X}_d + \lambda_s + \varepsilon_d. \quad (7)$$

By differencing outcomes and exposures between the initial and final periods, this specification removes all time-invariant district characteristics by construction. State fixed effects λ_s absorb state-level differences in long-run trends, policies, and institutional changes. Standard errors are clustered at the district level throughout.

6 Results

Section 6.1 first quantifies seasonal heterogeneity in climatic responses to upwind irrigation exposure. Section 6.2 studies how these climatic changes affect agricultural development.

6.1 Climate Spillovers

Opposing Seasonal Effects. Rainfall, temperature, and irrigation exhibit strong seasonal variations in India, shaped by the monsoon cycle, with a wet *Kharif* season (June–September) and a dry *Rabi* season (October–March), as explained in Section 2. Figure 4 illustrates the seasonal variation of irrigation on precipitation and temperature over the calendar year.

A 10 percentage-point increase in upwind irrigation exposure raises monthly precipitation by about 5 mm during the dry *Rabi* season. Relative to the sample mean of 26 mm per month during *Rabi*, this corresponds to an increase of about 19 percent. By contrast, upwind irrigation reduces rainfall by a similar magnitude during the wet *Kharif* season. This decline is concentrated in June, at the onset of the monsoon. The timing of the effect suggests that irrigation may delay the arrival of monsoon rainfall, a margin that is potentially important for agricultural production because delayed monsoon onset can affect planting decisions and crop development over the rest




Figure 4: Seasonal Effect of Upwind Irrigation on Climate

Notes: Figure 4 reports the estimated coefficients from specification (4) on (i) monthly precipitation (left axis, mm) and (ii) mean daily temperature (right axis, °C). Each estimated coefficient γ_m measures the *total* marginal effect of a 10 percentage point increase in upwind irrigation exposure in month m . Upwind irrigation is defined as the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance, averaged over the district. All regressions are estimated at the district level (corresponding to 311 districts) and include district fixed effects, month fixed effects and state-by-year fixed effects. Standard errors are clustered at the district level. Shaded areas represent 90% confidence intervals.

of the season.

Temperature effects mirror the precipitation response. Upwind irrigation lowers temperatures during the dry season, with an average cooling effect of about 0.3°C over *Rabi* and a peak of 0.5°C in January. In contrast, it increases temperatures during *Kharif*, with the largest warming effect again occurring in June.

Figure A13 in the Appendix shows that the same seasonal pattern extends to other climate indicators relevant for crop production. First, I examine the number of rainy days (precipitation above 1 mm), since higher total rainfall does not necessarily improve crop production if it is concentrated in a few intense rainfall events rather than distributed across the growing season

(Fishman, 2016). Consistent with the precipitation results, upwind irrigation increases rainy days during the *Rabi* season but reduces them in June, reinforcing the interpretation of a delayed monsoon onset. Second, I examine days with temperatures above 29°C, a threshold harmful for crop yields (Schlenker and Roberts, 2009). Upwind irrigation reduces hot days by about one during the dry season but increases them during the rainy season, with an additional seven hot days in June.

Matching Theory. Turning to the interpretation, these seasonal effects of irrigation on local climate are consistent with predictions from climate modeling studies (Guimberteau et al., 2012b; Agrawal et al., 2019; Mathur and AchutaRao, 2020). During the *Rabi* season, increased irrigation enhances evapotranspiration, promoting cloud formation, increasing precipitation, and lowering surface temperatures. However, this cooling also affects the subsequent monsoon. By reducing land-surface temperatures during the months preceding the monsoon, intensive upwind irrigation weakens the land–atmosphere thermal contrast that drives monsoon circulation, delaying its onset and reducing early monsoon precipitation. The empirical evidence presented in this paper closely matches these predictions and, to the best of my knowledge, provides the first empirical evidence supporting this mechanism. A direct test of the onset channel supports this reading: ten percentage points of upwind irrigation delay the start of the rainy season by 0.308 days (s.e. 0.114) and shorten it by 2.366 days (s.e. 0.460), against sample means of 151 and 341 days (Table A10, Appendix C.1).¹⁷

Taken together, these findings show that large-scale irrigation reshapes downwind weather, generating a climatic externality that depends critically on monsoon dynamics. As a result, the net effect on downwind agriculture is theoretically ambiguous. Better dry-season conditions may

¹⁷Onset and length are large-scale synoptic features, correlated across an entire state in a given year, so the state-by-year fixed effects used elsewhere in the paper would absorb most of their true variation and leave little for upwind exposure to explain. Table A10 instead controls for a state-specific linear trend in year, which still guards against divergent long-run climate trends across states; see Appendix C.1 for the comparison across fixed-effect structures. The contraction in length is robust across all of them; the delay in onset is significant only under this specification and is reported as suggestive.

raise land productivity and facilitate double cropping, while weaker monsoon conditions may depress yields during India’s main agricultural season. The model developed in Section 3 predicts that both climatic changes observed here strengthen incentives to adopt irrigation downstream: lower *Kharif* rainfall increases the value of irrigation as a substitute for rainfall, while higher *Rabi* rainfall raises the profitability of a second cropping season. Together, these mechanisms imply a self-reinforcing process whereby irrigation expansion upstream increases incentives to adopt irrigation downstream.

6.2 Downwind Development and Adaptation

Agricultural Outputs. Table 1 reports the estimated effects of upwind irrigation on downwind agricultural outcomes for the two major cereal crops in India: rice and wheat. As discussed in Section 2, rice is primarily cultivated during the monsoon season, while wheat is the dominant dry-season crop. Table 1 reports effects on the logarithm of total production, cultivated area, and yields.¹⁸ A 10 percentage point increase in upwind irrigation is associated with a 3.3% increase in wheat production. The rice coefficient has the same sign and a comparable magnitude, 2.7%, but is not distinguishable from zero on the unrestricted panel. For both crops, these gains are entirely driven by an expansion in cultivated area, which increases by 3.1% for wheat, again significant, and by 1.4% for rice, again not. Yield effects are small throughout: 0.3% for wheat and 1.3% for rice, the latter significant at the 5% level. These results indicate that the gains from upwind irrigation operate through the extensive margin of cultivation rather than through improvements in land productivity. At first glance, this pattern is puzzling, as upwind irrigation generates seasonally opposing climatic effects. Yet wheat production increases despite the absence of any detectable yield gain, and rice moves in the same direction. This suggests that the climatic spillovers generated by irrigation are mediated by farmers’ adaptation decisions rather than reflecting only the direct effects of weather on crop productivity.

¹⁸Yields are constructed as total cereal production divided by cultivated area.

Adaptation. I next examine how farmers adapt to the climatic changes induced by upwind irrigation. I focus on five potential margins of adjustment: irrigation, double cropping, mechanization, crop switching, and input use.

Table 2 shows that farmers primarily respond by expanding both irrigation and cropping intensity. A 10 percentage point increase in upwind irrigation raises cropping intensity, the ratio of gross to net cropped area, by 0.016 units (1.2% relative to its sample mean of 1.314), and irrigation intensity, the ratio of net irrigated to net cropped area, by 0.017 units (4.9% relative to its sample mean of 0.351). This expansion is accompanied by changes in irrigation sources (Table 3). The share of net irrigated area supplied by sources other than canals, tanks and wells rises by 0.9 percentage points, while the canal share declines by 1.2 percentage points. The tubewell share moves in the expected direction, by 0.6 percentage points, but is not significant on this panel. The negative coefficient on canal irrigation is reassuring, as it suggests that the results are not mechanically driven by spatial correlation with the upwind canal irrigation exposure variable. Instead, farmers appear to rely increasingly on privately controlled water sources. The growing importance of tubewells, which tap deeper aquifers and are typically privately managed, is consistent with the well-documented expansion of groundwater irrigation in India.

Beyond irrigation, farmers could also adapt by altering their production choices. Figure 6 reports effects for rice and for three other major cereal crops in India: pearl millet, maize, and sorghum. There is little evidence of systematic switching across cereal types, particularly from less water-intensive to more water-intensive crops. Instead, farmers continue to cultivate a diversified cereal portfolio.

Adaptation through input use also appears limited. Column (5) of Table 4 reports total NPK fertilizer intensity: the coefficient is 2.99 kilograms per hectare of net cropped area against a mean of 81.9, some 3.6 percent, and it is not significant (s.e. 2.29). There is little evidence of a broad increase in fertilizer application.

Similarly, there is little evidence that mechanization constitutes an important adaptation mar-

gin. Table 4 shows statistically insignificant effects on tractors and power tillers, while the number of electric pumps declines by approximately 870, significant at the 10% level. This decline may reflect substitution toward other irrigation technologies or constraints on electric pump adoption, such as limited electrification or unreliable power supply. In either case, mechanization does not appear to be a primary channel of adjustment.

Taken together, these findings suggest that farmers respond to irrigation-induced climatic changes mainly by expanding irrigation and increasing cropping intensity. By contrast, there is much weaker evidence of adaptation through crop switching, mechanization, or broad changes in input use.

Disentangling the Drivers of Irrigation Adoption. The above results show that farmers respond to irrigation-induced climatic changes by expanding irrigation downwind. I now provide suggestive evidence on the mechanism underlying this response. The theoretical framework points to two potential channels. First, lower *Kharif* rainfall and greater climate variability increase the value of irrigation as both a substitute for rainfall and a form of insurance, raising incentives to irrigate rice. Second, improved *Rabi* conditions increase the profitability of a second cropping season, potentially encouraging irrigation of wheat.

Table 5 provides evidence that is more consistent with the first mechanism. A 10 percentage point increase in upwind irrigation raises irrigated rice area by about 10% and irrigated wheat area by about 4%; the same exercise on four further crops shows the response concentrated in maize and sorghum, with no gain for pearl millet or sugarcane. Measured on dry-season rather than annual exposure — the season in which the induced rainfall actually falls — the wheat response is the cleaner of the two: 4.6% upwind against –1.3% on the reversed-wind arm. Although not definitive, this pattern suggests that downstream irrigation adoption serves both margins: offsetting weaker and more uncertain *Kharif* rainfall for rice, and extending the dry-season wheat crop where the induced rainfall makes a second season viable.

Farm Structure. While the previous results document how farmers are affected by and adapt to upwind irrigation exposure, they do not reveal who benefits from these adjustments. Adaptation may differ substantially across farms if investment capacity varies with farm size. To investigate this margin, I examine how exposure to upwind irrigation affects the farm size distribution. Results reported in Table 6 are estimated using five-year moving averages, as in specification (6).

In districts with greater exposure, the share of cultivated area operated by large farms (≥ 10 ha) rises by 0.75 percentage points (from a mean of 16.5%), largely at the expense of small (1–2 ha) and medium (4–10 ha) farms, whose shares decline by 0.35 and 0.39 percentage points, respectively. Importantly, the number of large farms does not increase significantly, indicating that this expansion reflects within-farm growth among existing large operations rather than the entry of new large farms. By contrast, marginal farms (< 1 ha) become more numerous, increasing by approximately 6,100 farms (significant at the 10% level), yet without meaningful growth in their operated area. This greater concentration of marginal holdings comes at the expense of medium-sized farms (4–10 ha), whose number declines by approximately 360 (significant at the 10% level).

These patterns point to a polarization of the agrarian structure in more exposed districts. Large farms expand their operated area, whereas smallholders—who constitute the vast majority of farms in India—remain trapped at the bottom of the farm-size distribution, likely because of limited investment capacity and lower adoption of productivity-enhancing technologies. As a result, the “missing middle” of medium-sized farms erodes, consistent with an incipient structural transformation in which irrigation-induced climate spillovers foster consolidation at the top and stagnation at the bottom.

7 Robustness

The identification strategy is subject to several potential threats, which I address in this section. Each exercise asks whether a plausible alternative source of variation could generate the esti-

mated climate and agricultural effects; the corresponding results are reported in Appendix D.

7.1 Placebo Specifications

Placebo Specifications. The identifying restriction is directional: irrigation should affect a district only when it lies on the air mass’s upwind path. Irrigation measured along the reversed, downwind path should therefore have no effect. I test this implication by estimating

$$y_{d,t} = \beta_1 E_{d,t}^{\uparrow} + \beta_2 E_{d,t}^{\downarrow} + \gamma \mathbf{X}_{d,t} + \mu_d + \lambda_{s,t} + \varepsilon_{d,t}, \quad (8)$$

where $E_{d,t}^{\downarrow}$ is exposure to irrigation located downwind of district d . Under the proposed mechanism, the upwind coefficient should remain stable while β_2 should be close to zero. This is what the data show in the dry season: the upwind estimates are stable and downwind coefficients are small and insignificant for both climate and agricultural outcomes (Appendix E). A similarly signed downwind effect would instead suggest broad spatial co-movement rather than atmospheric moisture transport. The placebo has one clear limitation. From July to September, monsoon circulation reverses, so the dry-season downwind arm overlaps mechanically with the monsoon upwind arm. The downwind coefficient is consequently not a valid placebo in these months and indeed fails there. I therefore treat the dry-season estimates as the cleanest evidence and interpret monsoon precipitation estimates more cautiously.

Geographic Confounders. The identifying assumption is that, conditional on the fixed effects, changes in upwind irrigation exposure are unrelated to other changes that affect outcomes. A concern is that districts more exposed to upwind irrigation may also have distinct geographic characteristics that generate different climate or agricultural trajectories. For example, if more exposed districts are flatter or more alluvial and those features shape their evolution over time, the estimates could attribute these differential trends to irrigation exposure. This concern is empirically relevant: highly exposed districts are flatter, lower, more alluvial, and more irrigated at baseline (Appendix Table A14). District and pixel fixed effects absorb such time-invariant differ-

ences, but not geography-specific trends. I therefore allow the eight principal geographic characteristics to have separate year effects, individually and jointly. The full set changes monthly climate coefficients by at most 0.80 mm and 0.005°C, and leaves agricultural estimates within the baseline confidence intervals (Appendix Figures A18–A19).

The identifying variation is concentrated in the Gangetic plain. If the results were driven by an unobserved shock or a feature unique to this region, the full-sample estimates could reflect that local factor rather than the proposed mechanism. I therefore re-estimate the model excluding each state, including Uttar Pradesh, and each of the 311 districts in turn. No estimate falls outside the full-sample confidence band. These exercises cannot rule out every unobserved geographic trend, but they show that neither observed geography nor a single location drives the results.

Wind-Borne Confounders. Wind is the source of identifying variation, but it can transport more than moisture. This creates a confounding risk when irrigated source regions covary with other wind-borne atmospheric inputs that affect climate or agriculture downwind. Crop-residue burning is the leading concern in India: irrigation can increase cropping intensity, burning releases aerosols, and those aerosols are carried downwind. The concern has the right signs to matter. Aerosols can cool the surface while altering the monsoon and precipitation (Bollasina et al., 2011; Rosenfeld et al., 2008; Ramanathan et al., 2005); rainfall also scavenges particulate matter, so the relationship between pollution and precipitation is bidirectional. Consistent with this mechanism, upwind exposure raises winter PM_{2.5} (Appendix Figure A24), showing that the identifying wind corridor carries particulate pollution alongside moisture and that this confounding risk must be assessed directly.

No single pollution control can resolve this issue. Contemporaneous PM_{2.5} begins only in 1998 and is partly affected by rainfall, so conditioning on it may introduce post-treatment bias. I therefore combine three tests. First, on the common post-2000 sample, controlling for PM_{2.5} barely changes the January temperature coefficient; a predetermined pollution measure interacted with year produces the same conclusion in the full panel. Second, I exclude the burn-

ing regions, the burning months and the burning era: the monthly climate profile is unchanged throughout (Appendix Figure A26), and rebuilding annual exposure with the residue-burning months taken out of the treatment leaves wheat production at 0.033 to 0.034 against 0.033 in the baseline (Appendix Table A16). Finally, I control for forest cover and fertiliser use. These variables capture other source-region changes correlated with initial irrigation expansion, including fertiliser-related emissions and afforestation, that may also affect the atmosphere. They leave the estimates qualitatively unchanged, although they too may respond to treatment. Taken together, the tests make pollution and related land-use changes unlikely explanations for the findings.

The Green Revolution. The expansion of irrigation projects overlapped with the Green Revolution. If districts most exposed to upwind irrigation were also those where high-yielding varieties, fertiliser, or other agricultural inputs expanded most rapidly, the estimates could capture differential agricultural modernization rather than irrigation-induced climate. State-by-year fixed effects absorb policies and input diffusion common to a state in a given year; the remaining concern is therefore differential adoption within a state and year.¹⁹ I first test this implication directly. Appendix Table A15 regresses each candidate confounder on exposure under the paper’s own fixed effects, in percent of the confounder’s residual standard deviation: none moves, the largest estimate being high-yielding varieties at 14.5 percent with a standard error of 9.2, and the reversed-wind placebo is null as well. I then add contemporaneous adoption measures and, separately, predetermined baseline measures interacted with year. The wheat coefficient remains positive, between 0.028 and 0.050 (Appendix Figure A23). These controls are imperfect, since contemporaneous adoption may itself respond to treatment and some series are incomplete, but the two exercises together weigh against differential Green Revolution trends as the explanation, without fully separating variables that trend together at the district level.

¹⁹District-specific linear trends remove the agricultural estimates, but they absorb 69% of the cumulative, nearly monotone variation in exposure. Such trends therefore leave little independent treatment variation with which to identify the effect.

Spatial Dependence and Functional Form. Weather observations are not independent draws. Rainfall and temperature are serially correlated within a district, and a shock at one location is likely to be shared by nearby pixels and districts. Clustering standard errors by district permits arbitrary serial correlation within districts, but treats districts as independent. I therefore report Conley standard errors that allow correlation to decay with distance over cutoffs from 100 to 1,000 km. Temperature estimates are unchanged; precision falls for August precipitation and wheat production, so these results should be read as less precisely determined than the corresponding temperature and wheat-area results (Appendix Figures A27–A28). I also estimate agricultural outcomes by PPML to retain zero observations, which account for 15% of district-years for wheat. The crop results strengthen and no longer depend on Uttar Pradesh. PPML preserves the dry-season climate findings; its pooled monsoon rainfall coefficient changes sign because additive and multiplicative models impose different cross-month restrictions when rainfall varies sharply over the year. Levels remain the preferred specification because the object of interest is a water balance measured in millimetres. A last concern is with inference itself rather than with the estimate: upwind exposure is smooth in space, smooth in time and almost monotonically increasing, and regressing such a variable on outcomes that are themselves smooth, trending and spatially correlated can produce apparently significant coefficients under the null. Reassigning each district’s entire exposure series to another district 500 times and re-estimating equation (4) on each draw, no permutation reaches the seasonal climate coefficients or the two intensity margins; wheat production clears 5 percent under Poisson and 10 percent in log, and rice cannot be told from a randomly assigned treatment. The share of permutations the paper’s own standard errors would call significant at 5 percent runs from 0.048 to 0.092, so the clustered errors are close to correctly sized and mildly optimistic where Poisson is used, and the randomisation p -value is the one to read there (Appendix D.8).

Canal Construction Dates and Canal Placement. Canal completion dates are sometimes measured imprecisely, and irrigation may become operational gradually rather than in the

recorded year. Such timing error could mechanically attenuate annual estimates or make them sensitive to short-run noise around construction. Re-estimating the model on five- and ten-year moving averages smooths both concerns. The dry-season effects remain, while the temperature effects persist and strengthen over longer horizons (Appendix Figure A14).

A related concern is that canals are not sited at random. Where a command area is built is set by the storage capacity of the reservoir and by the gradient of the land it can command (Section 4), and the qualitative record describes siting in those engineering terms — river flow, dam capacity, topography — and never with reference to the districts lying downwind. Endogenous placement is in any case a weaker threat here than in a study of a district’s own irrigation, since the treatment excludes the focal district’s own irrigated area and is identified by the rotation of the wind between the source region and the receiving district within the year. Exposed districts do differ in geography (Appendix Table A14); allowing each characteristic its own year effects leaves the estimates unchanged (Appendix Figures A18–A19).

8 Policy Implications

I use a stylized accounting exercise to quantify the effects of upwind irrigation on India’s dry-season climate and wheat sector. I first remove its estimated historical contribution, then project the effects of command areas under construction. This is not an equilibrium prediction: it holds prices, land, labour and capital fixed, and the identifying assumption is within-state-year, aggregated here as if it measured a level.

Historical importance of spillovers. I construct district-year counterfactuals by netting out the estimated effects of upwind irrigation. Between 1967 and 2020, dry-season rainfall declined by 10.9 mm; without upwind irrigation, the estimated decline would have been 27.0 mm. Temperature rose by 0.39°C over the same period, compared with a counterfactual increase of 0.76°C. Upwind irrigation therefore offset a substantial share of the drying and warming that would otherwise have occurred, with substantial geographical heterogeneity (see Appendix F, Figure A35).

The same exercise shows that these climatic spillovers account for 7 percent of the 1967–2010 increase in wheat production and 14 percent of the increase in wheat area. By 2006–2010, they generated 4.9 million tonnes of annual wheat production, worth 1.0 billion USD a year (19 USD per hectare of command area), and increased both irrigation and cropping intensity by 1.7 percentage points (Table A17, Panel A).

Future irrigation expansion. Canal command areas currently under construction would add 32.3 million hectares of large-scale irrigation and raising irrigation exposure by 10 percentage points. Applying my estimates implies 48 km³ of additional dry-season rainfall per year, equivalent to 8.4 percent of India’s observed dry-season rainfall, and 3.6 million additional tonnes of wheat, or 4.6 percent of sample production. Figure 5 maps the geographical distribution of these externalities with 30% of the benefits falling outside any command area (Table A17, Panel B).




Figure 5: Projected effects of canal irrigation schemes under construction

Notes: Both panels report the increment from the command areas recorded as ongoing or completed after 2010, a further 10.2 percentage points of area-weighted exposure. Panel (a) is the dry-season rainfall it would deliver, $\sum_m \hat{\beta}_m E_{i,m}$ summed over the Rabi months on the 0.25° grid and divided by the pixel's mean annual rainfall; Panel (b) is the share of district wheat production attributable to the same increment, on the 311 ICRISAT districts. The pixel grid is clipped to those districts, so the two panels cover the same territory and share one colour scale, trimmed at the 98th percentile of the two series pooled. Every command area under construction is outlined in red, at its actual size. Exposure is a leave-own-district-out average, so no unit is credited with a canal of its own; both panels show rainfall imported from elsewhere.

9 Conclusion

This paper provides new causal evidence that large-scale irrigation infrastructure can generate climatic spillovers by altering atmospheric dynamics and redistributing rainfall and temperature downwind. Leveraging exogenous variation in wind trajectories and five decades of high-resolution data on irrigation, wind, and rainfall, the empirical strategy isolates irrigation-induced climatic feedbacks from local agricultural decisions and shows that irrigation redistributes water across space and time with persistent consequences for downwind development trajectories.

The climatic effects are strongly seasonal: upwind irrigation increases rainfall during the dry Rabi season but reduces precipitation during the monsoon Kharif season and delays monsoon onset. A simple model shows how this seasonal rainfall reallocation can raise incentives to invest in downwind irrigation by increasing the value of dry-season cultivation while raising the marginal value of irrigation during the monsoon.

Consistent with this mechanism, irrigation-induced rainfall increases agricultural output primarily along the extensive margin through double cropping rather than yield gains. Farmers invest in private irrigation, particularly tubewells, generating a feedback loop whereby irrigation alters the climate and further encourages irrigation adoption. By contrast, farmers do not adapt through greater mechanization, crop switching, or increased use of yield-enhancing inputs. The benefits of these climatic spillovers are unevenly distributed: larger farms expand while smallholders remain stagnant, leading to increasing polarization in the farm size distribution and early signs of structural reallocation of land outside agriculture.

A further hint of this reallocation appears in the appendix: net cropped area falls and non-agricultural use rises with upwind exposure, alongside a weak decline in the female field wage (Tables A12 and A13). Whether this reflects an incipient structural transformation out of agriculture is a question this paper does not pursue and leaves to future research.

Are these climate effects realistic given the scale of the canal network? A simple water balance suggests they are. The regression implies 37 km³ ([18, 55]) of additional dry-season rainfall a year, or 8.9 percent ([4.2, 13.0]) of the region's observed total – an order-of-magnitude check, not

the network's absolute precipitation impact, since the coefficient is identified from relative differences across districts. On the supply side, the network irrigates some sixteen million hectares net, of which roughly ten are watered in the dry season; against the same land left unwatered, that is on the order of 30 km^3 of additional evapotranspiration each Rabi season, of which the recycling ratios typical of South Asia – around one fifth continentally, less over the hundred-kilometre ring identifying the estimates – would return anywhere from a few cubic kilometres to, at the most generous accounting, some 35 km^3 as local rainfall: roughly the same order as the regression estimate. An independent method reaches the same conclusion: using Lagrangian atmospheric moisture tracking rather than a wind-exposure regression, [Koppa et al. \(2026\)](#) estimate that upwind irrigation contributes 7.0 ± 5.2 percent of rainfall over India's rainfed cropland nationally, rising above 20 percent in the Indo-Gangetic Plain – both the magnitude and the geography match the results here.

Taken together, the findings highlight an underappreciated dimension of development infrastructure: its ability to generate climatic spillovers that reshape economic incentives and adaptation strategies. More broadly, the results underscore the importance of accounting for the indirect effects of infrastructure and the methodological challenges for causal identification when economic activity feeds back into the climate system.

Table 1: Effect of Upwind Irrigation on Rice and Wheat Outputs

	(1)	(2)	(3)	(4)	(5)	(6)
	<i>Rice – in log()</i>			<i>Wheat – in log()</i>		
	Production	Area	Yield	Production	Area	Yield
Upwind Irrigation (10 pp.)	0.027 (0.020)	0.014 (0.019)	0.013** (0.006)	0.033** (0.014)	0.031** (0.013)	0.002 (0.006)
Mean	4.063	3.826	0.221	3.731	3.368	0.332
Std. Dev.	2.246	2.028	0.634	2.295	2.072	0.575
Obs.	12,728	12,764	12,717	11,582	11,637	11,572
Nb. District	303	303	303	281	283	281
Within R^2	0.0012	0.0004	0.0010	0.0021	0.0023	0.0001
Fixed Effects	District and State $\times$ Year					

Notes: Each coefficient corresponds to the parameter β in equation (5), which captures the estimated effect of a 10 percentage point increase in upwind irrigation exposure. Production, cultivated area, and yields are reported separately for rice and wheat and expressed in log. Regressions include district fixed effects and state-by-year fixed effects, and no baseline controls: neither log population nor crop-specific log high-yielding variety area enters any column. Standard errors are clustered at the district level. Significance levels are denoted as follows: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 2: Effect of Upwind Irrigation on Cropping and Irrigation Intensities

	(1) Irrigation Intensity	(2) Cropping Intensity
Upwind Irrigation (10 pp.)	0.017*** (0.005)	0.016*** (0.005)
Mean	0.351	1.314
Std. Dev.	0.273	0.243
Obs.	12,472	13,300
Nb. District	311	311
Within R^2	0.0221	0.0099
Fixed Effects	District and State×Year	

Notes: Each coefficient corresponds to the parameter β in equation (5), which captures the estimated effect of a 10 percentage point increase in upwind irrigation exposure. Upwind irrigation is defined as the average share of irrigated land in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance. Irrigation intensity is measured as the ratio of net irrigated area (NIA) to net cropped area (NCA). Cropping intensity is measured as the ratio of gross cropped area (GCA) to net cropped area, and serves as a proxy for double cropping. Means and standard deviations are computed over the estimation sample. All regressions are estimated at the district level (corresponding to 311 districts based on 1966 administrative boundaries) and include district fixed effects as well as state-by-year fixed effects, without baseline controls. Standard errors, clustered at the district level, are reported in parentheses. Significance levels are denoted as follows: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 3: Effect of Upwind Irrigation on Irrigation Composition

	(1)	(2)	(3)	(4)	(5)
	<i>Dep. Var.: Share of Net Irrigated Area Attributed to ...</i>				
	Canals	Wells	Tubewells	Other Wells	Other Sources
Upwind Irrigation (10 pp.)	-0.012* (0.007)	0.005 (0.007)	0.006 (0.005)	-0.001 (0.007)	0.009*** (0.003)
Mean	0.273	0.506	0.210	0.297	0.114
Std. Dev.	0.293	0.341	0.288	0.320	0.235
Obs.	12,380	12,360	12,380	12,362	12,349
Nb. District	311	310	311	310	310
Within R^2	0.0033	0.0005	0.0012	0.0000	0.0026
Fixed Effects	District and State $\times$ Year				

Notes: Each coefficient corresponds to the parameter β in equation (5), which captures the estimated effect of a 10 percentage point increase in upwind irrigation exposure. Upwind irrigation is defined as the average share of irrigated land in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance. Results are reported for the shares of net irrigated area allocated to different irrigation practices. The sum of all shares in the sample equals one. Means and standard deviations are computed over the estimation sample. All regressions are estimated at the district level (corresponding to 311 districts based on 1966 administrative boundaries) and include district fixed effects as well as state-by-year fixed effects, without baseline controls. Standard errors, clustered at the district level, are reported in parentheses. Significance levels are denoted as follows: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.




Figure 6: Effect of Upwind Irrigation by Crop

Notes: Figure 6 reports the estimated coefficients from specification (5), the effect of a 10 percentage point increase in upwind irrigation exposure, for rice and for three other major cereals: maize, pearl millet and sorghum. Rice is the crop of the main text and is drawn here for comparison. Upwind irrigation is the mean share of irrigated land in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance. Area is in thousand hectares, production in tonnes, yield is production over area, and value is yield times the district’s annual crop price from ICRISAT. All four dependent variables are in logs, as marked, and estimated on the positive observations. Regressions include district and state-by-year fixed effects and *no baseline controls*: neither log population nor log high-yielding-variety area enters. Standard errors are clustered at the district level. Thick bars are 90 percent intervals, thin bars 95 percent; estimates not significant at 10 percent are drawn at reduced opacity.

Table 4: Effect of Upwind Irrigation on Mechanization and Fertilizer Use

	(1)	(2)	(3)	(4)	(5)
	Number of ... ('000)				t/ha of NCA
	Diesel Pump	Electric Pump	Power Tiller	Tractors	NPK
Upwind Irrigation (10 pp.)	0.057 (0.735)	-0.867* (0.451)	-0.075 (0.090)	-0.002 (0.251)	2.988 (2.290)
Mean	14.673	16.179	0.490	4.425	81.932
Std. Dev.	21.078	25.705	1.583	8.533	94.269
Obs.	1,903	1,913	1,511	2,013	13,070
Nb. District	311	311	311	311	311
Within R^2	0.0000	0.0018	0.0012	0.0000	0.0029
Specification	5-yr mean exposure, anchor years				Annual
Fixed Effects	District and State $\times$ Year				

Notes: Columns (1)–(4) report the parameter β of equation (6): upwind exposure averaged over the five years preceding the anchor year, the anchor year itself excluded, on the Livestock Census anchor years 1966, 1972, 1977, 1982, 1987, 1992, 1997, 2003 and 2007. Column (5) reports total NPK fertilizer per hectare of net cropped area on the annual panel under equation (5). This table is estimated without the baseline controls: neither log population nor log high-yielding-variety area enters any column. The HYV control is a sample restriction rather than a control — its coverage ends in 1989 in Uttar Pradesh and in 1987 in West Bengal, so including it deletes the two largest eastern rice states over the last two decades of the panel. Upwind irrigation is defined as the average share of irrigated land in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance. All regressions are estimated at the district level (311 districts on 1966 administrative boundaries) and include district and state-by-year fixed effects. Standard errors, clustered at the district level, are in parentheses. Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 5: Effect of Upwind Irrigation on Crop-Specific Irrigated Area

	(1)	(2)	(3)	(4)	(5)	(6)
	Irrigated Area – $\log()$					
	Rice	Wheat	Maize	Sorghum	Pearl millet	Sugarcane
Upwind Irrigation (10 pp.)	0.101** (0.050)	0.037* (0.020)	0.098** (0.049)	0.154** (0.062)	-0.013 (0.098)	-0.065* (0.038)
Mean	2.624	3.020	-0.140	-0.618	0.022	0.914
Std. Dev.	2.419	2.163	2.289	2.585	2.483	2.180
Obs.	10,679	10,522	8,235	5,023	4,450	10,316
Nb. District	297	277	264	211	173	285
Within R^2	0.0086	0.0024	0.0052	0.0058	0.0000	0.0034
Fixed Effects	District and State $\times$ Year					

Notes: Each coefficient corresponds to the parameter β in equation (5), which captures the estimated effect of a 10 percentage point increase in upwind irrigation exposure. Log irrigated area is reported for six crops, estimated on the positive observations. Regressions include district fixed effects and state-by-year fixed effects, and no baseline controls: neither log population nor crop-specific log high-yielding variety area enters any column. Standard errors are clustered at the district level. Significance levels are denoted as follows: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table 6: Effect of Upwind Irrigation on Farm-Size

	(1) Marginal (<1 ha)	(2) Small (1–2 ha)	(3) Semi-medium (2–4 ha)	(4) Medium (4–10 ha)	(5) Large (≥ 10 ha)	(6) Total
	<i>Dep. Var.: Share of Total Area</i>					<i>Area ('000 ha)</i>
Upwind Irrigation (10 pp.)	0.003 (0.003)	-0.003** (0.001)	-0.002 (0.003)	-0.004* (0.002)	0.007** (0.003)	3.977 (3.042)
Mean	0.167	0.180	0.229	0.259	0.165	535.047
Std. Dev.	0.152	0.084	0.068	0.112	0.161	333.874
Obs.	2,014	2,014	2,014	2,014	2,014	2,043
Nb. District	311	311	311	311	311	311
Within R^2	0.0024	0.0043	0.0011	0.0034	0.0075	0.0013
	<i>Dep. Var.: Number of Farms ('000)</i>					
Upwind Irrigation (10 pp.)	6.080* (3.131)	-0.535 (0.730)	-0.368 (0.464)	-0.361* (0.213)	0.118 (0.131)	4.901 (3.400)
Mean	171.926	58.045	41.375	24.562	6.296	304.031
Std. Dev.	184.954	40.515	26.894	21.128	9.625	214.213
Obs.	2,181	2,181	2,181	2,181	2,181	2,206
Nb. District	311	311	311	311	311	311
Within R^2	0.0040	0.0003	0.0003	0.0003	0.0006	0.0020
Fixed Effects	District and State $\times$ Year					

Notes: Each coefficient is the parameter β of equation (6): upwind exposure averaged over the five years *preceding* the anchor year, the anchor year itself excluded, with the dependent variable measured at the anchor year. Anchor years are those of the Agricultural Census: 1970, 1976, 1980, 1985, 1990, 1995, 2000 and 2005. Panel A reports shares of total farm area, built on the reported total and truncated to $[0, 1]$, and total area in column (6); Panel B reports the number of farms in each class. Regressions include district and state-by-anchor-year fixed effects and no baseline controls; this table carried none in the earlier draft either, so its estimates are unchanged by the removal of log population and log high-yielding-variety area. Upwind irrigation is the average share of irrigated land in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance. Standard errors, clustered at the district level, are in parentheses. Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

References

- AGRAWAL, S., A. CHAKRABORTY, N. KARMAKAR, ET AL. (2019): “Effects of winter and summer-time irrigation over Gangetic Plain on the mean and intra-seasonal variability of Indian summer monsoon,” *Climate Dynamics*, 53, 3147–3166.
- ALTER, R. E., E.-S. IM, AND E. A. B. ELTAHIR (2015): “Rainfall Consistently Enhanced around the Gezira Scheme in East Africa Due to Irrigation,” *Nature Geoscience*, 8, 763–767.
- ARAUJO, R. (2023): “When Clouds Go Dry: An Integrated Model of Deforestation, Rainfall, and Agriculture,” *Working Paper*.
- ASHER, S., A. CAMPION, D. GOLLIN, AND P. NOVOSAD (2022): “The Long-Run Development Impacts of Agricultural Productivity Gains: Evidence from Irrigation Canals in India,” CEPR Discussion Paper No. 17414.
- ASHER, S. AND P. NOVOSAD (2020): “Rural Roads and Local Economic Development,” *American Economic Review*, 110, 797–823.
- BASSINI, E., M. F. GIORDANO, AND R. E. NAMARA (2019): “The Impacts of Irrigation: A Review of Published Evidence,” Policy research working paper, World Bank Group, Washington, D.C., english version.
- BLAKESLEE, D., A. DAR, R. FISHMAN, S. MALIK, H. S. PELLEGRINA, AND K. S. BAGAVATHINATHAN (2023): “Irrigation and the Spatial Pattern of Local Economic Development in India,” *Journal of Development Economics*, 161, 102997.
- BLAKESLEE, D., R. FISHMAN, AND V. SRINIVASAN (2020): “Way Down in the Hole: Adaptation to Long-Term Water Loss in Rural India,” *American Economic Review*, 110, 200–224.
- BLANC, E. AND E. STROBL (2014): “Is Small Better? A Comparison of the Effect of Large and Small Dams on Cropland Productivity in South Africa,” *The World Bank Economic Review*, 28, 545–576.
- BOLLASINA, M. A., Y. MING, AND V. RAMASWAMY (2011): “Anthropogenic Aerosols and the Weakening of the South Asian Summer Monsoon,” *Science*, 334, 502–505.
- BOMBARDI, R. J., J. L. KINTER, AND O. W. FRAUENFELD (2019): “A Global Gridded Dataset of the Characteristics of the Rainy and Dry Seasons,” *Bulletin of the American Meteorological Society*, 101, E1315–E1328.

- BRAUN, T. AND W. SCHLENKER (2023): “Cooling Externality of Large-Scale Irrigation,” NBER Working Paper No. 30966.
- BURGESS, R., O. DESCHÊNES, D. DONALDSON, AND M. GREENSTONE (2017): “Weather, Climate Change and Death in India,” London School of Economics Working Paper.
- BURKE, M. AND K. EMERICK (2016): “Adaptation to Climate Change: Evidence from US Agriculture,” *American Economic Journal: Economic Policy*, 8, 106–140.
- BUSTOS, P., B. CAPRETTINI, AND J. PONTICELLI (2016): “Agricultural Productivity and Structural Transformation: Evidence from Brazil,” *American Economic Review*, 106, 1320–1365.
- BUSTOS, P., G. GARBER, AND J. PONTICELLI (2020): “Capital Accumulation and Structural Transformation*,” *The Quarterly Journal of Economics*, 135, 1037–1094.
- BUTLER, A. AND C. BOUDOT REDDY (2023): “Watering the Seeds of the Rural Economy: Evidence from Groundwater Irrigation in India,” *Available at SSRN 4575998*.
- BUTT, N., P. A. DE OLIVEIRA, AND M. H. COSTA (2011): “Evidence That Deforestation Affects the Onset of the Rainy Season in Rondonia, Brazil,” *Journal of Geophysical Research: Atmospheres*, 116.
- CARLETON, T. A. AND S. M. HSIANG (2016): “Social and Economic Impacts of Climate,” *Science*, 353, aad9837.
- CHOU, C., D. RYU, M.-H. LO, H.-W. WEY, AND H. M. MALANO (2018): “Irrigation-Induced Land–Atmosphere Feedbacks and Their Impacts on Indian Summer Monsoon,” *Journal of Climate*, 31, 8785–8801.
- CONLEY, T. G. (1999): “GMM Estimation with Cross Sectional Dependence,” *Journal of Econometrics*, 92, 1–45.
- CRUZ, J.-L. AND E. ROSSI-HANSBERG (2024): “The Economic Geography of Global Warming,” *Review of Economic Studies*, 91, 899–939.
- DEGOIS, M., T. VENDRYES, AND J. WOLFERSBERGER (2025): “Clearing the Way to Conflict: Deforestation and Downwind Rainfall,” *Working Paper*.
- DELL, M., B. F. JONES, AND B. A. OLKEN (2012): “Temperature Shocks and Economic Growth: Evidence from the Last Half Century,” *American Economic Journal: Macroeconomics*, 4, 66–95.
- (2014): “What Do We Learn from the Weather? The New Climate-Economy Literature,” *Journal of Economic Literature*, 52, 740–98.

- DEVANAND, A., M. HUANG, M. ASHFAQ, B. BARIK, AND S. GHOSH (2019): “Choice of Irrigation Water Management Practice Affects Indian Summer Monsoon Rainfall and Its Extremes,” *Geophysical Research Letters*, 46, 9126–9135.
- DILLON, A. AND R. FISHMAN (2019): “Dams: Effects of Hydrological Infrastructure on Development,” *Annual Review of Resource Economics*, 11, 125–148.
- DONALDSON, D. (2018): “Railroads of the Raj: Estimating the Impact of Transportation Infrastructure,” *American Economic Review*, 108, 899–934.
- DONALDSON, D. AND R. HORNBECK (2016): “Railroads and American Economic Growth: A “Market Access” Approach,” *The Quarterly Journal of Economics*, 131, 799–858.
- DUFLO, E. AND R. PANDE (2007): “Dams*,” *The Quarterly Journal of Economics*, 122, 601–646.
- EMERICK, K. (2018): “Agricultural Productivity and the Sectoral Reallocation of Labor in Rural India,” *Journal of Development Economics*, 135, 488–503.
- FAO (2023): “Land Use Statistics and Indicators 2000–2021: Global, Regional and Country Trends,” FAO-STAT Analytical Briefs Series 71, FAO, Rome.
- FEDER, G., R. JUST, AND D. ZILBERMAN (1985): “Adoption of Agricultural Innovations in Developing Countries: A Survey,” *Economic Development and Cultural Change*, 33, 255–98.
- FICK, S. E. AND R. J. HIJMANS (2017): “WorldClim 2: New 1-km Spatial Resolution Climate Surfaces for Global Land Areas,” *International Journal of Climatology*, 37, 4302–4315.
- FILLON, R. (2025): “The Biophysical Channels of Climate Impacts,” *Working Paper*.
- FISHMAN, R. (2016): “More uneven distributions overturn benefits of higher precipitation for crop yields,” *Environmental Research Letters*, 11, 024004.
- FISHMAN, R., M. JAIN, AND A. KISHORE (2013): “Groundwater Depletion, Adaptation and Migration: Evidence from Gujarat, India,” IFPRI Discussion Paper.
- FOSTER, A. D. AND M. R. ROSENZWEIG (1995): “Learning by Doing and Learning from Others: Human Capital and Technical Change in Agriculture,” *Journal of Political Economy*, 103, 1176–1209.
- GREVE, P., A. U. SCHMITT, D. G. MIRALLES, S. McDERMID, K. L. FINDELL, A. GARCÍA-GARCÍA, AND J. PENG (2025): “Observational Evidence of Increased Afternoon Rainfall Downwind of Irrigated Areas,” *Nature Communications*, 16, 3415.

- GROSSET-TOUBA, F., A. PAPP, AND C. TAYLOR (2024): “Rain Follows the Forest: Land Use Policy, Climate Change, and Adaptation,” Available at SSRN: <https://ssrn.com/abstract=4333147>.
- GUIMBERTEAU, M., K. LAVAL, A. PERRIER, AND J. POLCHER (2012a): “Global Effect of Irrigation and Its Impact on the Onset of the Indian Summer Monsoon,” *Climate Dynamics*, 39, 1329–1348.
- (2012b): “Global effect of irrigation and its impact on the onset of the Indian summer monsoon,” *Climate Dynamics*, 39, 1329–1348.
- HANSEN, Z. K., G. D. LIBECAP, AND S. E. LOWE (2011): “Climate Variability and Water Infrastructure: Historical Experience in the Western United States,” in *The Economics of Climate Change: Adaptations Past and Present*, National Bureau of Economic Research, Inc., NBER Chapters, 253–280.
- HENDERSON, J. V., A. STOREYGARD, AND U. DEICHMANN (2017): “Has Climate Change Driven Urbanization in Africa?” *Journal of Development Economics*, 124, 60–82.
- HORNBECK, R. AND P. KESKIN (2014): “The Historically Evolving Impact of the Ogallala Aquifer: Agricultural Adaptation to Groundwater and Drought,” *American Economic Journal: Applied Economics*, 6, 190–219.
- (2015): “Does Agriculture Generate Local Economic Spillovers? Short-Run and Long-Run Evidence from the Ogallala Aquifer,” *American Economic Journal: Economic Policy*, 7, 192–213.
- HROZENCIK, R. A., M. WOERMAN, L. PREONAS, AND F. BURLIG (2025): “Evidence of Climate Adaptation: Temperature and Groundwater Extraction in California Agriculture,” *Working Paper*.
- HUSSAIN, I. AND M. A. HANJRA (2004): “Irrigation and poverty alleviation: review of the empirical evidence,” *Irrigation and Drainage*, 53, 1–15.
- IMD (2026): “Frequently Asked Questions on Monsoon,” <https://mausam.imd.gov.in/imd/latest/monsoonfaq.pdf>, accessed June 2026.
- JACK, B. K. AND N. RYAN (2026): “Economic Development and the Environment,” NBER Working Paper number 34735.
- JONES, M., F. KONDYLIS, J. LOESER, AND J. MAGRUDER (2022): “Factor Market Failures and the Adoption of Irrigation in Rwanda,” *American Economic Review*, 112, 2316–52.
- KALA, N. (2017): “Learning, Adaptation, and Climate Uncertainty: Evidence from Indian Agriculture,” MIT CEEPR Working Paper Series 2017-023.

- KAPA, H., K. BHARGHAVI, AND T. L. REDDY (2025): “Variability in Indian monsoon onset: Delays, advances, and regional disruptions,” *Evolving Earth*.
- KOPPA, A., F. BASSANI, J. KEUNE, V. DEMAN, D. INSUA-COSTA, V. HARI, S. GHOSH, D. G. MIRALLES, AND S. BONETTI (2026): “Irrigated Agriculture Supports Rainfed Crops in India through Atmospheric Moisture Recycling,” *Environmental Research Letters*, 21.
- KRUSELL, P. AND A. A. J. SMITH (2022): “Climate Change Around the World,” NBER Working Paper no. 30338.
- LEITE-FILHO, A. T., B. S. SOARES-FILHO, J. L. DAVIS, G. M. ABRAHÃO, AND J. BÖRNER (2021): “Deforestation Reduces Rainfall and Agricultural Revenues in the Brazilian Amazon,” *Nature Communications*, 12, 2591.
- LI, T. (2021): “Protecting the Breadbasket with Trees? The Effect of the Great Plains Shelterbelt Project on Agriculture,” *Land Economics*, 97, 321–344.
- LIU, M., Y. SHAMDASANI, AND V. TARAZ (2023): “Climate Change and Labor Reallocation: Evidence from Six Decades of the Indian Census,” *American Economic Journal: Economic Policy*, 15, 395–423.
- LOARIE, S. R., D. B. LOBELL, G. P. ASNER, Q. MU, AND C. B. FIELD (2011): “Direct Impacts on Local Climate of Sugar-Cane Expansion in Brazil,” *Nature Climate Change*, 1, 105–109.
- LOBELL, D. B., W. SCHLENKER, AND J. COSTA-ROBERTS (2011): “Climate Trends and Global Crop Production Since 1980,” *Science*, 333, 616–620.
- MANOLI, G., S. FATICHI, M. SCHLÄPFER, K. YU, T. W. CROWTHER, N. MEILI, P. BURLANDO, G. G. KATUL, AND E. BOU-ZEID (2019): “Magnitude of Urban Heat Islands Largely Explained by Climate and Population,” *Nature*, 573, 55–60, pMID: 31485056.
- MATHUR, R. AND K. A. ACHUTARAO (2020): “A modelling exploration of the sensitivity of India’s climate to irrigation,” *Climate Dynamics*, 54, 1851–1872.
- NOCCO, M. A., R. A. SMAIL, AND C. J. KUCHARIK (2019): “Observation of irrigation-induced climate change in the Midwest United States,” *Global Change Biology*, 25, 3472–3484.
- OLMSTEAD, S. M. AND H. SIGMAN (2015): “Damming the Commons: An Empirical Analysis of International Cooperation and Conflict in Dam Location,” *Journal of the Association of Environmental and Resource Economists*, 2, 497–526.
- PAI, D. S., L. SRIDHAR, M. RAJEEVAN, O. P. SREEJITH, N. S. SATBHAI, AND B. MUKHOPADHYAY (2014): “De-

- velopment of a new high spatial resolution ($0.25^\circ \times 0.25^\circ$) long period (1901–2010) daily gridded rainfall data set over India and its comparison with existing data sets over the region,” *MAUSAM*, 65, 1–18, India Meteorological Department, Pune.
- PUMA, M. J. AND B. I. COOK (2010): “Effects of irrigation on global climate during the 20th century,” *Journal of Geophysical Research*, 115, D16120.
- RAMANATHAN, V., C. CHUNG, D. KIM, T. BETTGE, L. BUJA, J. T. KIEHL, W. M. WASHINGTON, Q. FU, D. R. SIKKA, AND M. WILD (2005): “Atmospheric Brown Clouds: Impacts on South Asian Climate and Hydrological Cycle,” *Proceedings of the National Academy of Sciences*, 102, 5326–5333.
- RILEY, S. J., S. D. DEGLORIA, AND R. ELLIOT (1999): “A Terrain Ruggedness Index That Quantifies Topographic Heterogeneity,” *Intermountain Journal of Sciences*, 5, 23–27.
- ROSENFELD, D., U. LOHMANN, G. B. RAGA, C. D. O’DOWD, M. KULMALA, S. FUZZI, A. REISELL, AND M. O. ANDREAE (2008): “Flood or Drought: How Do Aerosols Affect Precipitation?” *Science*, 321, 1309–1313.
- RYAN, N. AND A. SUDARSHAN (2020): “Rationing the Commons,” NBER Working Paper no. w27749.
- SAEED, F., S. HAGEMANN, AND D. JACOB (2009): “Impact of Irrigation on the South Asian Summer Monsoon,” *Geophysical Research Letters*, 36, L20711.
- SCHLENKER, W. AND M. J. ROBERTS (2009): “Nonlinear Temperature Effects Indicate Severe Damages to U.S. Crop Yields under Climate Change,” *Proceedings of the National Academy of Sciences*, 106, 15594–15598.
- SEKHRI, S. (2014): “Wells, Water, and Welfare: The Impact of Access to Groundwater on Rural Poverty and Conflict,” *American Economic Journal: Applied Economics*, 6, 76–102.
- SINGH, O., A. KASANA, AND P. BHARDWAJ (2022): “Understanding energy and groundwater irrigation nexus for sustainability over a highly irrigated ecosystem of north western India,” *Applied Water Science*, 12, 44.
- SPRACKLEN, D. V., S. R. ARNOLD, AND C. M. TAYLOR (2012): “Observations of Increased Tropical Rainfall Preceded by Air Passage over Forests,” *Nature*, 489, 282–285.
- STROBL, E. AND R. O. STROBL (2011): “The Distributional Impact of Large Dams: Evidence from Cropland Productivity in Africa,” *Journal of Development Economics*, 96, 432–450.
- TARAZ, V. (2017): “Adaptation to Climate Change: Historical Evidence from the Indian Monsoon,” *Environment and Development Economics*, 22, 517–545.

- (2018): “Can Farmers Adapt to Higher Temperatures? Evidence from India,” *World Development*, 112, 205–219.
- TAYLOR, C. A. (2025): “Irrigation and Climate Change: Long-run Adaptation and Its Externalities,” *Working Paper*.
- TE WIERIK, S. A., E. L. H. CAMMERAAT, J. GUPTA, AND Y. A. ARTZY-RANDRUP (2021): “Reviewing the Impact of Land Use and Land-Use Change on Moisture Recycling and Precipitation Patterns,” *Water Resources Research*, 57, e2020WR029234, e2020WR029234 2020WR029234.
- URFELS, A., A. J. McDONALD, T. J. KRUPNIK, AND P. R. VAN OEL (2020): “Drivers of groundwater utilization in water-limited rice production systems in Nepal,” *Water International*, 45, 39–59.
- VAN DONKELAAR, A., M. S. HAMMER, L. BINDLE, M. BRAUER, J. R. BROOK, M. J. GARAY, N. C. HSU, O. V. KALASHNIKOVA, R. A. KAHN, C. LEE, R. C. LEVY, A. LYAPUSTIN, A. M. SAYER, AND R. V. MARTIN (2021): “Monthly Global Estimates of Fine Particulate Matter and Their Uncertainty,” *Environmental Science & Technology*, 55, 15287–15300.
- WADA, Y., L. P. H. VAN BEEK, C. M. VAN KEMPEN, J. W. T. M. RECKMAN, S. VASAK, AND M. F. P. BIERKENS (2010): “Global depletion of groundwater resources,” *Geophysical Research Letters*, 37.
- ZAVERI, E., J. RUSS, AND R. DAMANIA (2020): “Rainfall anomalies are a significant driver of cropland expansion,” *Proceedings of the National Academy of Sciences*, 117, 10225–10233.

Appendix

A. Data and descriptive statistics

- Descriptive statistics — Table A7.
- The seasonal cycle of precipitation and temperature — Figure A7.
- Precipitation and temperature by season — Figure A8.
- Temperature distribution by season across India — Figure A9.
- The prevailing wind by season — Figure A10.
- Canal irrigation potential — Table A8.
- Upwind irrigation exposure, the level and the variation — Figure A11.

B. Model

- Reduced-form adoption frontier and comparative statics — Figure A12.

C. Additional climate results

- Monthly climate response, district and pixel level — Figure A13.
- Monthly climate response under three estimation horizons — Figure A14.
- Net annual effect on precipitation — Table A9.
- Monsoon onset and length, state-trend specification — Table A10.
- Non-linear effect across the temperature distribution — Figure A15, Table A11.

D. Robustness

- Sensitivity to the fixed-effect structure — Figure A16.
- Weighting and distance band of the exposure — Figure A17.
- District characteristics regressed on upwind exposure — Table A14.
- Monthly climate coefficients at pixel level — Figure A18.
- Agricultural coefficients under each geographic control — Figure A19.
- Leave-one-state-out, monthly climate — Figure A20.
- Leave-one-state-out, agricultural outcomes — Figure A21.
- Log, Poisson and levels on the agricultural outcomes — Figure A22.
- Green Revolution: upwind exposure and the candidate confounders — Table A15, Figure A23.
- Air pollution and crop-residue burning — Figures A24 to A26, Table A16.
- Conley spatial standard errors, 100 to 500 km — Figures A27 and A28.
- Randomisation inference, the treatment reassigned 500 times — Figure A29.

E. Placebo — the wind-connection matrix reversed

- Wheat production, area, yield and irrigated area — Figure A30.
- Cropping and irrigation intensity — Figure A31.

F. Counterfactual exercise

- The counterfactual, standing network and projects under construction — Table A17.
- State by state, the projects under construction — Figure A33.
- State by state, the network standing by 2010 — Figure A34.
- Where the counterfactual leaves the identifying range — Figure A35.

A Data and Descriptive Statistics

Table A7: Descriptive statistics of the estimation sample

	Obs.	Mean	Standard deviation		p10	Median	p90	% zero
			Total	Within				
Treatment								
Upwind irrigation exposure (%), district-month	205,161	17.9	26.0	11.7	0.11	5.42	62.5	7.0
Upwind irrigation exposure (%), district-year	13,683	16.6	24.0	7.62	0.41	5.61	56.0	2.5
Climate								
Precipitation (mm/month)	205,161	95.0	147	138	0.00	26.8	288	10.9
Rabi, October–March	102,621	26.1	53.7	48.6	0.00	4.81	76.4	19.1
Kharif, June–September	68,383	226	179	134	47.7	189	440	–
Mean daily temperature (°C)	205,161	24.3	6.53	4.77	16.8	25.4	31.2	–
Rabi, October–March	102,621	20.9	6.23	3.22	15.1	22.2	26.5	–
Kharif, June–September	68,383	26.9	4.10	1.83	23.5	27.1	31.2	–
Degree days above 29°C (days/month)	205,161	18.3	39.5	37.6	0.00	0.00	73.5	47.2
Rabi, October–March	102,621	1.20	5.05	4.82	0.00	0.00	1.66	76.3
Kharif, June–September	68,383	20.2	39.3	34.2	0.00	1.07	76.9	25.2
Agriculture								
Wheat production (000 t)	13,566	158	305	145	0.00	40.0	445	14.6
Wheat area (000 ha)	13,578	73.7	95.6	29.2	0.00	35.7	199	14.2
Rice production (000 t)	13,613	205	295	138	0.40	89.8	538	6.5
Rice area (000 ha)	13,614	128	159	31.9	0.60	65.1	350	6.2
Cropping intensity (GCA/NCA)	13,300	1.31	0.24	0.13	1.05	1.27	1.64	–
Irrigation intensity (GIA/GCA)	12,355	0.34	0.26	0.10	0.05	0.27	0.76	0.1
Adaptation margins								
Wheat irrigated area (000 ha)	12,420	62.1	93.5	35.5	0.00	19.0	183	15.2
Rice irrigated area (000 ha)	12,413	57.6	94.0	33.7	0.00	12.1	173	13.9
NPK fertiliser per ha of net cropped area	13,070	81.9	94.3	69.2	4.32	49.1	209	0.7
Tractors (000)	2,292	3.97	8.11	5.85	0.03	0.95	10.8	4.1
Electric pumps (000)	2,202	15.2	25.3	15.5	0.03	4.93	42.5	5.7

Notes: Statistics are computed on the estimation sample. The climate rows are district-months over 1966–2020; every other row is district-years over 1966–2009, on the 311 districts of the 1966 boundaries. *Total* is the standard deviation over all observations, *within* the standard deviation once each district mean is removed — the variation the district fixed effects leave to identify β . Climate variables are reported pooled and then by season; *Rabi* is October to March and *Kharif* June to September. The last column is the share of observations equal to zero, a dash meaning fewer than one in two thousand.




Figure A7: The seasonal cycle of precipitation and temperature

Notes: Figure A7 plots the seasonal cycle of monthly precipitation and mean daily temperature over 1966–2020, on the 311 districts of the estimation sample. Each district’s own average for each calendar month is computed first; the line is the median of those normals across districts and the bands their 25th–75th and 5th–95th percentile ranges. Shaded columns mark Kharif, June to September, and Rabi, October to March.




Figure A8: Precipitation and temperature by season

Notes: Figure A8 maps the 1966–2020 normals of rainfall and mean daily temperature by season, on the 4,244 pixels of the 0.25° grid the exposure measure is built on. Rabi is October to March and Kharif June to September. The two rainfall panels carry their own legend, since the monsoon delivers an order of magnitude more water than the dry season; the two temperature panels share one scale in 2°C classes, so a colour denotes the same temperature in both, and classes a season never reaches are dropped from its legend.




Figure A9: Temperature distribution by season across India

Notes: Figure A9 plots the distribution of mean daily temperature across district-month normals, 1966–2020, one density per cropping season. Each district’s own average for each calendar month is computed first, so the unit is a district-month normal. Rabi is October to March and Kharif June to September.




Figure A10: The prevailing wind by season


Notes: Figure A10 is drawn from the backward-trajectory matrices at 800 hPa on which the exposure measure of Section 5.1 is built, averaged over 1966–2010. A link joins a pixel to a cell that sends it air; links are kept when their length falls in the 50–150 km window the treatment uses, and each arrow is the mean displacement over the links retained, averaged within one-degree cells and coloured by its length in degrees. Rabi is October to March and Kharif June to September.

Table A8: Canal Irrigation Potential of India

	Area (ha)	% of Completed	% of India
Completed (known year)	56,249,134	96.500	17.210
Completed (unknown year)	2,037,783	3.500	0.620
Total Completed	58,286,917	100.000	17.830
Ongoing	27,233,781		8.330
Total Potential	85,520,698		26.160

Notes: The table reports the total area (in hectares) under canal irrigation according to WRIS in India. “Completed (known year)” refers to projects with an identified year of completion, while “Completed (unknown year)” includes those for which the completion year is unavailable. “Ongoing” refers to projects under construction during the observation period. % of Completed indicates the share of completed irrigation potential relative to the total completed projects, whereas % of India indicates the share relative to the total surface of the country.

A. The level, 1966–2010 average



B. The share of the identifying variation, 1966–2010




Figure A11: Upwind irrigation exposure: the level and the variation

Notes: Both panels are drawn on the 1966 district boundaries, the 311 districts of the estimation sample, and report upwind irrigation exposure: the share of irrigated area in pixels lying 50–150 km upwind of the district, weighted inversely by distance, with irrigation inside the district itself excluded. Panel A is its 1966–2010 mean, in percent. Panel B is the within-district standard deviation of the same series over the same window, in percentage points — the variation the district and state-by-year fixed effects leave to identify β .

B Reduced-Form Representation and Closed-Form Frontier

To provide additional analytical insight into the comparative statics of the model, I derive a reduced-form representation of the irrigation adoption frontier. The objective is not to impose additional structure on the production functions but rather to approximate the two margins identified in the main text while preserving their theoretical properties.

Recall that the gain from irrigation adoption can be decomposed into two components: a substitution effect in the monsoon season and an activation effect in the dry season. To obtain a tractable expression, I approximate each component with reduced-form functions that capture the qualitative restrictions implied by the model.

First, I represent the activation margin as linear in effective water availability during the Rabi season:

$$A(r_R) = ar_R, \quad a > 0,$$

where r_R denotes water availability during the dry season and a captures the productivity of activated Rabi cultivation.

Second, the substitution gain in the Kharif season is represented as decreasing in monsoon rainfall:

$$S_K(r_K) = \frac{b}{r_K + \varepsilon}, \quad b > 0.$$

This functional form captures the idea that irrigation is most valuable when monsoon rainfall is scarce. As rainfall increases, the buffering value of irrigation declines. The expression can be interpreted as a reduced-form approximation of the difference $f_K(r_K + \alpha_K) - f_K(r_K)$ under concavity of the production function. Under these approximations, the adoption gain can be written as

$$\Delta(r_K, r_R) = ar_R + \frac{b}{r_K + \varepsilon} - \bar{F}, \quad (\text{A9})$$

where $\bar{F}$ denotes the fixed cost of installing irrigation.

Closed-Form Adoption Frontier The irrigation adoption frontier is defined by the set of rainfall realizations for which the farmer is indifferent between adopting irrigation and not adopting

it: $\Delta(W_K, W_R) = 0$. Solving for the level of dry-season water availability that satisfies this condition yields the frontier

$$r_R^*(r_K) = \frac{\bar{F}}{a} - \frac{b}{a(r_K + \varepsilon)}. \quad (\text{A10})$$

This equation characterizes the combinations of seasonal rainfall (r_K, r_R) for which irrigation adoption yields zero net gain.

Interpretation The frontier partitions the (r_K, r_R) plane into two regions: (i) If $r_R > r_R^*(r_K)$, then $\Delta > 0$ and irrigation is adopted. (ii) If $r_R < r_R^*(r_K)$, irrigation is not adopted.

The slope of the frontier is

$$\frac{dr_R^*}{dr_K} = \frac{b}{a(r_K + \varepsilon)^2} > 0. \quad (\text{A11})$$

Higher monsoon rainfall therefore raises the level of dry-season rainfall required for irrigation to be profitable. Intuitively, when Kharif rainfall is abundant, the substitution value of irrigation is small, and stronger gains from dry-season production are needed to justify the fixed cost.

The curvature of the frontier is given by

$$\frac{d^2 r_R^*}{dr_K^2} = -\frac{2b}{a(r_K + \varepsilon)^3} < 0. \quad (\text{A12})$$

The frontier is therefore concave in r_K . This curvature reflects diminishing marginal substitution value in the monsoon season: as rainfall increases, the incremental buffering value of irrigation declines at a decreasing rate. Figure A12 illustrates the irrigation adoption frontier implied by the model and depicts its shifts according to the frontier's sensitivity to its key parameters.




Figure A12: Irrigation Adoption Frontier and Comparative Statics

Notes: The figure illustrates the irrigation adoption frontier implied by the reduced-form representation of the model. The horizontal axis shows water availability during the monsoon (Kharif) season r_K , while the vertical axis shows the minimum level of dry-season (Rabi) water $r_R^*(r_K)$ required for irrigation adoption. Irrigation is adopted whenever the pair (r_K, r_R) lies above the frontier, i.e., in the blue shaded area in Panel (1). Panel (2) shows the effect of higher Rabi productivity (a), which lowers the threshold for adoption by increasing the profitability of dry-season cultivation. Panel (3) shows the effect of a higher Kharif scarcity value (b), which increases the substitution value of irrigation when monsoon rainfall is low. Panel (4) illustrates the effect of higher fixed costs (F), which shifts the adoption frontier upward and reduces irrigation adoption for all rainfall realizations.

C Additional Results




Figure A13: Seasonal Effect of Upwind Irrigation on Climate

Notes: Figure A13 reports the estimated coefficients from specification (4) on (a) monthly precipitation (in mm), (b) number of rainy days (> 1mm), (c) mean daily temperature ($^{\circ}\text{C}$) and (d) number of harmful degree days (days > 29°C). Each estimated coefficient γ_m measures the *total* marginal effect of a 10 percentage point increase in upwind irrigation exposure in month m . Upwind irrigation is defined as the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance. Each panel carries the two levels of aggregation. The district regressions cover the 311 districts and include district, month and state-by-year fixed effects; the pixel regressions are on the 0.25° grid inside those districts and include pixel, month and district-by-year fixed effects. Standard errors are clustered at the district level in both. Shaded areas represent 90% confidence intervals.

C.1 Net Annual Impact of Irrigation on Climate

Specification. The analysis is conducted at the monthly level using observations available at the 0.25° degree grid and estimated as follows:




Figure A14: The monthly climate response under three estimation horizons

Notes: Figure A14 reports the estimated coefficients from specification (4) on monthly precipitation (mm) and mean daily temperature (°C), upwind arm alone, under four estimation horizons: the annual panel, five- and ten-year moving averages of each district-month cell, and a long difference between the two ends of the panel. Each estimated coefficient γ_m measures the *total* marginal effect of a 10 percentage point increase in upwind irrigation exposure in month m . Upwind irrigation is the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance, averaged over the district. The moving averages take five- and ten-year means of each district-month cell, with state-by-centre fixed effects; averaging is over years within a district-month cell and never across months. The long difference compares the 1966–1975 and 2011–2020 decade means of each district-month cell, with state and month fixed effects; a narrower window centred on the two ends of the panel produced noisy, wrongly-signed outliers and is not used. All four are estimated at the district level with standard errors clustered at the district level. Every line is solid; colour carries the horizon rather than significance.

$$Y_{i,d,t,m} = \beta e_{i,d,t,m} + \mu_i + \lambda_{d,t,m} + \varepsilon_{i,d,t,m}, \quad (\text{A13})$$

where $Y_{i,d,t,m}$ is the outcome of interest—climate variables—for pixel i in district d , year t , and month m . The key explanatory variable, *Upwind Irrigation* $e_{i,d,t,m}$, measures the share of upwind irrigated land affecting pixel i at time (t, m) . Pixel fixed effects, μ_i , absorb time-invariant location characteristics, while district-by-year-by-month fixed effects, $\lambda_{d,t,m}$, capture local shocks and seasonal variation. Since $e_{i,d,t,m}$ is defined as the proportion of upwind area under irrigation (ranging from 0 to 1), the coefficient β can be interpreted as the change in climate associated with moving from rainfed to fully irrigated upwind; it is reported in the paper’s own units, per 10 percentage points, in Table A9.

Results. Table A9 reports the net annual effect of upwind irrigation on downwind precipitation and temperature: a single coefficient with no month interaction, so it averages the opposing Rabi and Kharif responses documented in Section 6.1. The table is estimated twice, once on the pixel-month grid and once on the district-month panel, each clustered at its own level of observation. At pixel level a 10 percentage point increase in exposure raises monthly rainfall by 1.125 mm with a standard error of 0.367, adds 0.052 rainy days (0.017), raises mean daily temperature by 0.021°C (0.010) and removes 0.123 degree days above 29°C (0.060). Splitting the temperature distribution is informative: the whole of that mean warming sits in the cold tail. The coldest day of the month rises by 0.037°C (0.012), significant at one percent, while the hottest day moves by 0.008°C (0.008) and is indistinguishable from zero. The district panel reproduces the pattern, with a minimum-temperature coefficient of 0.083 (0.047) significant at ten percent and a maximum-temperature coefficient of 0.017 (0.030) that is not. Upwind irrigation therefore compresses the diurnal range from below rather than adding heat at the top: it lifts night-time and cold-season minima, which is the signature of the added atmospheric moisture that suppresses radiative cooling, and it leaves the hot extreme, the margin that matters for crop damage, untouched. That is consistent with the negative coefficient on degree days above 29°C in column (4). The district panel returns point

estimates of the same sign and larger magnitude, but none is distinguishable from zero. The two panels are not in conflict on the coefficients, only on their precision: clustering by district allows for the spatial correlation of upwind exposure across neighbouring pixels and roughly doubles every standard error, which is enough to remove significance. Read with that caveat, the net annual effect is small: a specification without a month interaction nets the Rabi gain against the Kharif loss and recovers little. It is reported here for completeness and is not used elsewhere in the paper.

Table A10 reports two further outcomes under the same equation, the start and duration of the rainy season, on their own specification. Onset and length are large-scale synoptic features rather than pixel-level ones, strongly correlated across a state in a given year; saturating the fixed effects with a state-by-year interaction, as in Table A9, absorbs almost all of that true regional variation and leaves β to identify off residual local noise. Replacing the interaction with a state-specific linear trend in year – which still controls for divergent long-run climate trends across states – raises ten percentage points of upwind irrigation to a 2.366-day contraction of the rainy season (0.460, $p < 0.01$), against 1.873 days under the fully saturated specification, and surfaces a 0.308-day delay in its onset (0.114, $p < 0.01$) that is null under every other fixed-effect structure tried. The length result is robust to this choice of specification; the onset result, while it survives excluding any one state at a time, is not, and is reported as suggestive rather than established.

This is also where the results depart from the closest reference. Grosset-Touba et al. (2024) study the 1930s Great Plains Shelterbelt and find that upwind forest cover raised precipitation and lowered temperature downwind for decades, over distances of up to 200 km, with downwind crop yields 11 to 22 percent higher. Their climate effect is a level shift: rainfall rises over the year, and the agricultural response follows from a uniformly more favourable growing season. The dry-season estimates here have the same sign and a comparable spatial reach – Rabi rainfall rises and Rabi temperature falls, over an upwind band of 50 to 150 km – but they do not survive aggregation to the year, because the monsoon response has the opposite sign. Irrigation in India redistributes rainfall across the calendar rather than adding to it, which is also why the agricultural response

documented here runs through cropping intensity and private irrigation rather than through yields.

Table A9: Effect of Upwind Irrigation on Climate: Net Annual Effect

	(1)	(2)	(3)	(4)
	Precipitation		Temperature	
	Sum (mm)	Rainy Days (Nb.)	Mean (°C)	29°C DD (Nb.)
Panel A: Pixel level				
Upwind irrigation (10 pp.)	1.125*** (0.367)	0.052*** (0.017)	0.021** (0.010)	-0.123** (0.060)
Mean	88.14	6.37	24.63	19.99
Obs.	2,262,653	2,262,653	2,257,571	2,257,571
Nb. Pixel	4,275	4,275	4,260	4,260
Fixed Effects	Pixel, Month and District×Year			
Panel B: District level				
Upwind irrigation (10 pp.)	1.639 (1.398)	0.092 (0.062)	0.047 (0.035)	-0.078 (0.157)
Mean	94.97	6.67	24.31	18.37
Obs.	167,857	167,857	167,857	167,857
Nb. District	311	311	311	311
Fixed Effects	District, Month and State×Year			

Notes: Each coefficient is the parameter β of equation (A13), estimated without a month interaction, so it is the net annual effect of a 10 percentage point increase in upwind exposure, averaging the opposing Rabi and Kharif responses of Figure 4. Precipitation outcomes are total monthly rainfall and the number of days above 1 mm; temperature is the monthly mean of daily temperature and degree days above 29°C. Panel A is the pixel-month panel on the 0.25° grid over 1966–2010; Panel B is the district-month panel over the same years. Estimated on upwind exposure without exclusion of the focal district’s own pixels (pixel level) or with it (district level, the paper’s own variable); standard errors are clustered at the level of observation of each panel – by pixel in Panel A, by district in Panel B. Standard errors are in parentheses. Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C.2 Non-linear Effect of Irrigation on Temperature

An important literature shows that agricultural output responds nonlinearly to temperature, with yields increasing at moderate temperatures but declining sharply beyond crop-specific thresholds (Schlenker and Roberts, 2009). Therefore, I extend equation (5) to examine whether irrigation exposure affects the distribution of temperature exposure.

Table A10: Effect of Upwind Irrigation on Monsoon Onset and Length: State-Trend Specification

	(1) Onset (days)	(2) Length (days)
Upwind Irrigation (10 pp.)	0.308*** (0.114)	-2.366*** (0.460)
Mean	151.080	341.260
Std. Dev.	21.431	138.334
Obs.	9,018	9,018
Nb. District	311	311
Within R^2	0.0003	0.0027
Fixed Effects	District, Year and State-specific Linear Trend	

Notes: Onset and Length are the start and duration of the rainy season, in days from January 1 (Onset) and in days (Length), from the pentad (5-day period) rainy/dry season detection of [Bombardi et al. \(2019\)](#) applied to CHIRPS precipitation, multiplied by five. A higher Onset is a later start to the rainy season. Each project pixel is the mean of the source dataset's 0.05° cells within 0.16° of its centre, averaged to the district-year panel of Table A9's Panel B, over 1981–2009. Onset and length are large-scale synoptic features, correlated across a state in a given year: the state-by-year fixed effects of Table A9 would absorb most of their true variation, so this table replaces that interaction with a state-specific linear trend in year, which still controls for divergent long-run climate trends across states. Standard errors are clustered at the district level. The Length coefficient is robust across every fixed-effect structure tried, from the state-by-year interaction of Table A9 to no state control at all; the Onset result is significant only under the specification reported here and is read as suggestive. Season length should be read with caution outside the core monsoon belt, where the retreating (northeast) monsoon can chain onto the southwest monsoon in the pentad detection, spanning most of the year. Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

To this end, I adopt a binned Growing Degree Days (GDD) approach, which measures cumulative heat exposure relevant for crop growth by aggregating daily temperatures into predefined temperature bins over the growing season, following [Schlenker and Roberts \(2009\)](#). Specifically, the temperature distribution is partitioned into a set of 1-degree bins, and for each bin a separate GDD variable is constructed. Each GDD variable aggregates, over all days of the growing season, the amount of heat received when daily temperatures fall within the corresponding temperature range. As a result, seasonal GDD capture cumulative thermal exposure in degree-days ($^\circ\text{C} \times \text{days}$), reflecting both the frequency and the intensity of heat within each temperature bin.²⁰

²⁰To illustrate, consider a temperature bin spanning $34\text{--}35^\circ\text{C}$ during a given growing season. If, over the season, temperatures reach 34.2°C on two days and 36°C on one day, the corresponding GDD in this bin would equal $0.2 + 0.2 + 1 = 1.4$ degree-days. In contrast, a season with a single day at 34.1°C would generate only 0.1 degree-day in this bin. Thus, seasonal GDD summarize overall exposure to heat within a temperature range, rather than simply

As previously, this analysis is conducted separately for the two main agricultural seasons in India, the Rabi and Kharif seasons, which differ in crop composition and climatic conditions. I estimate the following specification:

$$\text{GDD}_{d,s,t,b} = \sum_{k \in \mathcal{B}} \beta_k (E_{d,t} \times \mathbf{1}\{b = k\}) + \mu_d + \lambda_{s,t} + \delta_b + \varepsilon_{d,s,t,b}. \quad (\text{A14})$$

Here, $\text{GDD}_{d,s,t,b}$ denotes cumulative growing degree days accumulated in temperature bin b in district d , state s , and year t over an entire season, belonging to either the Rabi or the Kharif agricultural season. As in equation (5), $E_{d,t}$ corresponds to district-level upwind irrigation exposure, measured separately for the Rabi and Kharif seasons. The indicator $\mathbf{1}\{b = k\}$ allows irrigation exposure to have a fully flexible, nonlinear effect across the temperature distribution.

The coefficients β_k capture the marginal effect of irrigation exposure on cumulative degree days in temperature bin k , that is, on the total amount of heat accumulated within that temperature range over the season. The specification includes district fixed effects μ_d , state-by-year fixed effects $\lambda_{s,t}$, and temperature-bin fixed effects δ_b , which absorb time-invariant district characteristics, common state-level shocks in a given year, and average differences across temperature bins common to all districts, respectively. Standard errors are clustered at the district level.

Figure A15 illustrates the nonlinear relationship between upwind irrigation exposure and the distribution of temperature exposure across growing degree day bins. During the dry Rabi season, higher exposure to upwind irrigation reallocates thermal exposure away from higher temperature bins toward more moderate ones: cumulative degree days rise between 4 and 17°C and fall above 17°C, up to about 40°C. During the wet Kharif season the pattern is reversed. Cumulative degree days fall over the moderate range, most sharply between 20 and 25°C, and rise from 26°C upward. Irrigation therefore alters local thermal conditions in season-specific ways, mitigating extreme heat exposure during the dry season while amplifying high-temperature exposure during the monsoon.

counting the number of hot days.




Figure A15: Nonlinear Effect of Upwind Irrigation on Temperature Exposure by Season

Notes: Figure A15 plots the estimated coefficients $\hat{\beta}_b$ from equation (A14), which measure the marginal effect of upwind irrigation exposure on cumulative growing degree days in temperature bin b . Upwind irrigation exposure is defined as the mean share of irrigated area in pixels located 50–150 km upwind of the focal district, weighted inversely by distance. All regressions are estimated at the district level and include district fixed effects as well as state-by-year fixed effects. Standard errors are clustered at the district level. Shaded areas represent 95% confidence intervals. Estimates are reported separately for the Kharif (wet) season, shown in blue, and the Rabi (dry) season, shown in yellow.

Table A11: Effect of Upwind Irrigation on Climate Outcomes

	(1)	(2)	(3)	(4)	(5)	(6)
	<i>Climate outcomes</i>					
Season	Precipitation (mm/month)			Mean Temperature (°C)		
	Coef.	Mean(Y)	Effect (% Mean)	Coef.	Mean(Y)	Effect (% Mean)
Rabi	16.409*** (5.254)	29.09	56.41%	-0.276*** (0.066)	19.62	-1.40%
Zaid	4.178 (3.303)	50.49	8.27%	-0.024 (0.099)	27.71	-0.09%
Kharif	-7.226 (4.740)	215.14	-3.36%	0.073 (0.076)	25.83	0.28%
Obs.		768,938			766,932	
Fixed Effects		Pixel, and District $\times$ Season $\times$ Year				

Notes: Each coefficient is the season-specific marginal effect of upwind irrigation exposure moving from 0 (non-irrigated) to 1 (fully irrigated), estimated in one regression per outcome with a season-specific slope on exposure, pixel and district-by-season-year fixed effects. The outcome is the average monthly value within the season. Seasons are *Kharif* = June–September, *Zaid* = April–May, and *Rabi* = the six remaining months, October through March. Standard errors, clustered at the district level, are in parentheses; the same estimates appear as the last rung of Figure A16. Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

C.3 Land composition and field wages

Table A12: Effect of Upwind Irrigation on Land Composition

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
	<i>Dep. Var.: Share of total geographic area</i>							
	Non-agri.	Perm. past.	Net crop. area	Cult. waste	Cur. fallow	Oth. fallow	Barren	Forest
Upwind irrigation (10 pp.)	0.0014** (0.0007)	-0.0001 (0.0002)	-0.0027** (0.0012)	-0.0005 (0.0007)	0.0016* (0.0010)	0.0005 (0.0004)	-0.0007 (0.0005)	0.0006 (0.0011)
Mean	0.081	0.041	0.528	0.037	0.043	0.028	0.055	0.163
Std. Dev.	0.059	0.066	0.205	0.058	0.046	0.032	0.071	0.170
Obs.	13,357	13,356	13,325	13,357	13,356	13,354	13,357	13,354
Nb. District	311	311	311	311	311	311	311	311
Adj. R^2	0.908	0.928	0.962	0.907	0.771	0.841	0.893	0.956
Fixed Effects	District and State $\times$ Year							

Notes: Each coefficient is β of equation (5), the effect of a 10 percentage point increase in upwind irrigation exposure. Upwind irrigation is the average share of irrigated land in pixels 50–150 km upwind of the district, weighted by the inverse of distance and excluding the district's own area. Each outcome is a land-use category as a share of total geographic area; the eight categories together cover 99.4 percent of that area at the median, so the table is close to a complete decomposition. All columns include district and state-by-year fixed effects, and no baseline controls: neither log population nor high-yielding variety area enters any column. Standard errors are clustered at the district level in parentheses. Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

Table A13: Effect of Upwind Irrigation on Field Labour Wages

	(1)	(2)
	<i>Daily wage – in log()</i>	
	Male wage	Female wage
Upwind irrigation (10 pp.)	0.0043 (0.0077)	-0.0195* (0.0108)
Mean (log)	2.446	2.399
Std. Dev. (log)	1.239	1.229
Obs.	9,363	6,148
Nb. District	308	271
Adj. R^2	0.978	0.969
Fixed Effects	District and State $\times$ Year	

Notes: Each coefficient is β of equation (5), the effect of a 10 percentage point increase in upwind irrigation exposure. The dependent variables are the logarithm of the daily field labour wage for men and for women, available over 1966–2009. Both columns include district and state-by-year fixed effects, and no baseline controls, with standard errors clustered at the district level in parentheses. The female sample is markedly smaller than the male one, 271 districts against 308, because female field wages are reported less often. Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

D Robustness

D.1 Sensitivity to the fixed-effect structure




Figure A16: Sensitivity of the main results to the fixed-effect structure: a ladder

Notes: Each panel re-estimates the same regression, adding one fixed effect at a time, from none to the structure of the paper's own tables. The climate estimates are pooled across seasons with a season-specific slope on exposure and a season fixed effect at every rung, becoming district-by-season-by-year at the last; the agricultural estimates use the annual district panel. Upwind irrigation is the average share of irrigated land in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance. Standard errors are clustered at the district level. Error bars denote 90 and 95 percent confidence intervals. The last climate rung is the specification of Table A11 and reports the same coefficients.

D.2 The definition of exposure




Figure A17: Different Definitions of Upwind Irrigation

Notes: Figure A17 re-estimates the main results under sixteen alternative definitions of upwind exposure, varying the distance band (horizontal axis) and the weighting: circles weight by inverse distance, as the paper does, triangles weight equally. Panel A reports the seasonal coefficient of specification (4); Panel B reports β of equation (5). District and state-by-year fixed effects; standard errors clustered at the district level. Thick bars are 90% and thin bars 95% confidence intervals; the dotted line is the published coefficient.

D.3 Geographic confounders




Figure A18: Monthly climate coefficients at pixel level, under every geographic control

Notes: Figure A18 reports the monthly coefficient of a 10 percentage point increase in upwind irrigation exposure on precipitation (mm) and mean temperature (°C), estimated at the pixel level. The black line is the paper's specification, with no geographic control; each coloured line adds one time-invariant geographic characteristic interacted with year, and the last adds all of them together. Pixel, month and district-by-year fixed effects; standard errors clustered at the district level.

Table A14: Balance Table: District Characteristics on Upwind Irrigation Exposure

	Mean	S.D.	No controls		State FE		State FE + other X	
	(1)	(2)	$\hat{\beta}$ (3)	p (4)	$\hat{\beta}$ (5)	p (6)	$\hat{\beta}$ (7)	p (8)
Panel A: exposure in levels (mean 1966–2010)								
Latitude (°)	22.73	5	0.516***	0.000	0.062*	0.053	-0.059**	0.012
Longitude (°)	78.91	4.349	-0.010	0.830	-0.075*	0.091	-0.152***	0.001
Elevation (m)	297	212.8	-0.322***	0.000	-0.346***	0.001	-0.133**	0.023
Ruggedness (TRI)	12.74	15.55	-0.352***	0.000	-0.406***	0.003	0.007	0.238
Elevation s.d. (m)	88.34	95.58	-0.362***	0.000	-0.500***	0.000	-0.136***	0.009
Slope (°)	0.7728	0.9602	-0.346***	0.000	-0.418***	0.003	-0.005	0.404
Distance to coast (km)	440.7	307	0.667***	0.000	0.153***	0.001	0.150***	0.000
District area (km ²)	965.8	598.1	-0.381***	0.000	-0.179**	0.013	-0.088	0.335
Alluvial aquifer share	0.4289	0.4279	0.646***	0.000	0.336***	0.000	0.138**	0.017
Hard-rock aquifer share	0.1173	0.2107	-0.251***	0.000	-0.213**	0.012	0.048	0.595
Forest share, 1966–75	0.1514	0.1633	-0.351***	0.000	-0.332***	0.001	0.068	0.306
Irrigated share, 1966–75	0.2419	0.2109	0.673***	0.000	0.517***	0.000	-0.047*	0.088
Length of growing period (days)	188.7	48.21	-0.003	0.961	-0.122	0.150	0.021	0.805
Normal annual rainfall (mm)	1078	486.5	-0.179***	0.000	-0.191***	0.009	-0.051	0.122
Potential evapotranspiration	1567	184	-0.251***	0.000	0.034	0.616	-0.060	0.359
Moisture availability index	0.5661	0.3322	-0.124***	0.004	-0.203***	0.007	0.044	0.129
Rainfall 1968–72 (mm)	998	523	-0.230***	0.000	-0.185***	0.004	-0.062	0.223
Cropping intensity, 1968–72	1.19	0.1537	0.582***	0.000	0.246***	0.006	0.076	0.399
Irrigation intensity, 1968–72	0.2344	0.2038	0.695***	0.000	0.534***	0.000	0.070**	0.011
Imbalanced at 10%			17 / 19		17 / 19		8 / 19	
Joint F (geography → exposure)			38.57		5.82			
R^2			0.777		0.364 (within)			
Panel B: exposure in changes (1998–2002 minus 1968–1972)								
Latitude (°)	22.73	5	0.240***	0.000	0.010	0.626	-0.021*	0.051
Longitude (°)	78.91	4.349	0.263***	0.000	0.053*	0.055	0.018	0.540
Elevation (m)	297	212.8	-0.330***	0.000	-0.193***	0.000	-0.117***	0.000
Ruggedness (TRI)	12.74	15.55	-0.245***	0.000	-0.156***	0.005	0.004	0.111
Elevation s.d. (m)	88.34	95.58	-0.266***	0.000	-0.177***	0.002	-0.020	0.123
Slope (°)	0.7728	0.9602	-0.244***	0.000	-0.160***	0.005	-0.003	0.200
Distance to coast (km)	440.7	307	0.232***	0.000	0.037	0.220	0.027	0.180
District area (km ²)	965.8	598.1	-0.232***	0.000	-0.113***	0.006	-0.086**	0.047
Alluvial aquifer share	0.4289	0.4279	0.387***	0.000	0.166***	0.001	0.082***	0.004
Hard-rock aquifer share	0.1173	0.2107	-0.162***	0.000	-0.040	0.287	0.057	0.199
Forest share, 1966–75	0.1514	0.1633	-0.221***	0.000	-0.169***	0.001	-0.048	0.120
Irrigated share, 1966–75	0.2419	0.2109	0.087*	0.055	-0.020	0.668	-0.004	0.672
Length of growing period (days)	188.7	48.21	0.228***	0.000	0.025	0.491	0.064*	0.062
Normal annual rainfall (mm)	1078	486.5	-0.019	0.651	-0.095**	0.011	-0.028**	0.024
Potential evapotranspiration	1567	184	-0.193***	0.000	0.014	0.708	-0.032	0.294
Moisture availability index	0.5661	0.3322	0.014	0.746	-0.099**	0.015	0.016	0.182
Rainfall 1968–72 (mm)	998	523	-0.014	0.717	-0.090**	0.012	-0.033	0.240
Cropping intensity, 1968–72	1.19	0.1537	0.206**	0.000	0.008	0.899	-0.031	0.537
Irrigation intensity, 1968–72	0.2344	0.2038	0.034	0.443	-0.022	0.610	-0.001	0.926
Imbalanced at 10%			15 / 19		11 / 19		6 / 19	
Joint F (geography → exposure)			4.06		2.95			
R^2			0.335		0.184 (within)			

Notes: Cross-sectional balance on 272 districts in 13 states, one observation per district. Each row is a separate regression of the characteristic on upwind irrigation exposure; both are standardized, so $\hat{\beta}$ is the standard-deviation change in the characteristic per standard-deviation of exposure. Columns (3)–(4) have no controls, (5)–(6) add state fixed effects, (7)–(8) add every other characteristic in the table. Standard errors are heteroskedasticity-robust. The joint F tests all characteristics together in a regression of exposure on the full vector.




Figure A19: Agricultural outcomes under each geographic control

Notes: Figure A19 reports the estimated coefficients from $y_{d,t} = \beta E_{d,t}^{\uparrow} + \gamma(X_d \times \text{Year}_t) + \mu_d + \lambda_{s,t} + \varepsilon_{d,t}$ on four agricultural outcomes, adding one time-invariant geographic characteristic interacted with year (each row's label) and then all of them together. Each estimated coefficient β measures the marginal effect of a 10 percentage point increase in upwind irrigation exposure. Upwind irrigation is defined as the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance, averaged over the district. All regressions include district and state-by-year fixed effects and standard errors are clustered at the district level. Error bars denote 90% and 95% confidence intervals. The top row of each panel, in black, is the specification of the paper, without geographic controls; the rows below add one characteristic at a time and then all of them together. Wheat and rice production are in logs, as marked, and estimated on the positive observations; double cropping and irrigation intensity are bounded ratios and enter in levels, as in the main text.




Figure A20: Leave-one-state-out: monthly climate response

Notes: Figure A20 reports the estimated coefficients from specification (4) on (i) monthly precipitation (mm) and (ii) mean daily temperature ($^{\circ}\text{C}$), re-estimated fifteen times, each dropping one state from the sample. Each estimated coefficient γ_m measures the *total* marginal effect of a 10 percentage point increase in upwind irrigation exposure in month m . Upwind irrigation is defined as the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance, averaged over the district. All regressions are estimated at the district level and include district, month and state-by-year fixed effects. Standard errors are clustered at the district level. The black line is the full sample and the shaded area its 90% confidence interval. Each grey line is one leave-one-out sample, that is, the estimate obtained when one of the fifteen states is removed; the two whose removal displaces the coefficients furthest are drawn in colour and named in the legend. Panel B reports that displacement, in full-sample standard errors. Because the state-by-year fixed effects are separable, dropping a state does not alter how the others are identified.




Figure A21: Leave-one-state-out: agricultural outcomes

Notes: Figure A21 reports the estimated coefficient on upwind exposure from $y_{d,t} = \beta E_{d,t}^{\uparrow} + \mu_d + \lambda_{s,t} + \varepsilon_{d,t}$, re-estimated sixteen times, once on the full sample and once dropping each of the fifteen sample states. Each coefficient is the effect of a 10 percentage point increase in upwind irrigation exposure. All specifications include district and state-by-year fixed effects and no baseline controls, with standard errors clustered at the district level, so the top row reproduces Tables 1 to 2 exactly. Both estimators are drawn: OLS on $\log y$, on the positive observations, and Poisson pseudo-maximum likelihood on every observation; the two bounded ratios are least squares in levels under either heading and are drawn once. The top row is the full sample, the second isolates Uttar Pradesh with its own 90 and 95 percent intervals, and the bottom row is one point per remaining state with a segment spanning their range.

Uttar Pradesh and the crop results. Uttar Pradesh contributes 54 of the 311 sample districts, 35 percent of wheat production and 60 percent of the within-district variation in upwind exposure: it is both where the wheat is grown and where the specification is identified. Dropping it removes the wheat result in logs and produces a rice result that was not there before, which looks like a single state carrying the crop effects. It is not; it is the logarithm. Outside Uttar Pradesh 16.9 percent of district-year wheat observations are true zeros, against 3.3 percent inside, and $\log y$ discards every one of them, so the log specification is silent precisely where wheat is grown intermittently. Poisson pseudo-maximum likelihood, which keeps the zeros, reverses the conclusion

(Figure A21): dropping Uttar Pradesh *raises* the wheat coefficients rather than removing them, with wheat production at 0.070 ($p = 0.03$) outside against 0.030 ($p < 0.01$) on the full sample, and wheat area at 0.050 against 0.035. The rice response outside Uttar Pradesh, finally, is on area and not on yield — rice area 0.039 under Poisson and 0.092 in logs, against 0.015 and 0.014 on the full sample — which is acreage expansion under irrigation, the adaptation margin the paper describes, and not a direct effect of monsoon rainfall that the estimates report as falling.

D.4 Functional form




Figure A22: The main agricultural results under three functional forms

Notes: Figure A22 reports the estimated coefficient on upwind exposure from $y_{d,t} = \beta E_{d,t}^{\uparrow} + \mu_d + \lambda_{s,t} + \varepsilon_{d,t}$ under three functional forms: *log* (the main text's specification), *Poisson*, and *levels* divided by the outcome's mean. District and state-by-year fixed effects; standard errors clustered at the district level. Thick bars are 90% and thin bars 95% confidence intervals; estimates not significant at 10 percent are drawn at reduced opacity.

D.5 The Green Revolution

Table A15: Effect of Upwind Irrigation on Green Revolution Confounders

	(1)	(2)
	Effect	(s.e.)
Forest share of district land	-0.5%	(3.1)
Barren share of district land	-1.4%	(2.3)
NPK fertiliser per ha (NCA)	4.9%	(7.4)
HYV share of gross cropped area	14.5%	(9.2)
Log population	6.8%	(5.5)
Fixed Effects	District and State $\times$ Year	

Notes: Each row regresses a candidate confounder, measured at the focal district, on upwind exposure under the paper's own fixed effects. The test asks whether the treatment predicts the confounder at the receiving district. None does, at any conventional significance level. Coefficients are expressed as a percentage of the outcome's residual standard deviation for a 10 percentage point increase in exposure, since a coefficient is only small relative to the dispersion it is meant to explain. Standard errors are clustered at the district level. The same logic is applied to PM_{2.5} in Appendix D.6 and to varietal adoption and fertiliser use in Appendix D.5.




Figure A23: Agricultural results under Green Revolution controls

Notes: Figure A23 reports the estimated coefficient β from $y_{d,t} = \beta E_{d,t}^{\uparrow} + \mu_d + \lambda_{s,t} + \varepsilon_{d,t}$ on four agricultural outcomes, with the stated Green Revolution control added. Upwind irrigation is the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance, averaged over the district. The two markers are the share of gross cropped area under high-yielding varieties (HYV) and NPK fertiliser use. The panel is in three groups, separated by a grey rule, because each rests on a different sample: the top three rows – the paper’s own specification and the two predetermined baseline controls, each the marker’s 1966–1970 district mean interacted with year – use the full sample; the next two restrict to the district-years where HYV is observed, the paper’s specification on that subsample and then with contemporaneous HYV added; the last two do the same for NPK. Each point carries its own observation count, since the three groups differ in N ; comparisons should be read within a group, not across groups. Wheat and rice production are in logs, on the positive observations; double cropping and irrigation intensity are bounded ratios and enter in levels, as in the main text. Standard errors are clustered at the district level; error bars denote 90% and 95% confidence intervals.

D.6 Air pollution and crop-residue burning




Figure A24: Upwind irrigation exposure and PM_{2.5}

Notes: Figure A24 reports the estimated coefficients from specification (4) with PM_{2.5} as the dependent variable, 2000–2020. Each coefficient γ_m measures the *total* marginal effect of a 10 percentage point increase in upwind irrigation exposure in month m , in micrograms per cubic metre. Upwind irrigation is the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance. PM_{2.5} is the ACAG satellite product, extracted monthly on the paper’s own 0.25° grid. The district series is solid, with district, month and state-by-year fixed effects; the pixel series is dashed, with pixel, month and district-by-year fixed effects. Standard errors are clustered at the district level in both, since pollution is correlated far beyond a single 0.25° pixel. Black is all of India; orange restricts to the 89 districts of Punjab, Haryana, Uttar Pradesh and Bihar, where crop-residue burning is concentrated. The shaded columns are the two burning windows, April–May after the wheat harvest and October–November after the rice harvest. Shaded areas are 90% confidence intervals.




Figure A25: Climate Results controlling for PM 2.5

Notes: Estimated at the pixel level, with pixel, month and district-by-year fixed effects, clustered by district. Figure A25 reports the estimated coefficients from specification (4) on (i) monthly precipitation (mm) and (ii) mean daily temperature ($^{\circ}\text{C}$). *Main* is the paper's specification on the full 1966–2020 panel. *Post 2000* restricts it to the years $\text{PM}_{2.5}$ is observed, and *Post 2000 + PM2.5* adds the contemporaneous, continuous, pixel-level measure directly. Only the contemporaneous control is reported: a predetermined control, each pixel's mean $\text{PM}_{2.5}$ interacted with year, is constant within a pixel-year and can therefore only translate the twelve monthly coefficients as a block, never deform the seasonal profile the paper reads. Shaded areas are 90% confidence intervals.




Figure A26: Climate Results dropping the burning regions and months

Notes: Figure A26 reports the estimated coefficients from specification (4) on (i) monthly precipitation (mm) and (ii) mean daily temperature ($^{\circ}\text{C}$), re-estimated after removing the residue-burning regions, the burning months and the burning era. Each coefficient γ_m measures the *total* marginal effect of a 10 percentage point increase in upwind irrigation exposure in month m . Upwind irrigation is the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance, averaged over the district. All regressions are estimated at the district level with district, month and state-by-year fixed effects and standard errors clustered at the district level. *Main* is the full sample, drawn in black with its 90% confidence band; each coloured line drops the stated regions, months or years, and breaks over the months it no longer estimates.

Table A16: Effect of Upwind Irrigation on Agricultural Outcomes: Exclusions and Additional Controls

	(1)	(2)	(3)	(4)
	Wheat prod. (log)	Rice prod. (log)	Double cropping	Irrigation int.
Main	0.033** (0.014)	0.027 (0.020)	0.016*** (0.005)	0.017*** (0.005)
<i>Panel A. Burning months removed from the exposure measure</i>				
All months (rebuilt)	0.033** (0.014)	0.025 (0.020)	0.015*** (0.005)	0.014*** (0.004)
Excl. Oct–Nov (rice residue)	0.033** (0.014)	0.028 (0.019)	0.015*** (0.005)	0.014*** (0.004)
Excl. Apr–May (wheat residue)	0.034** (0.014)	0.025 (0.020)	0.015*** (0.005)	0.014*** (0.004)
<i>Panel B. Controlling for other Wind-Borne Confounders</i>				
No control (panel B sample)	0.035** (0.014)	0.026 (0.021)	0.015*** (0.005)	0.014*** (0.005)
+ forest share	0.033** (0.014)	0.027 (0.021)	0.016*** (0.005)	0.014*** (0.005)
+ fertiliser (NPK) per ha	0.029** (0.014)	0.019 (0.021)	0.012** (0.005)	0.012*** (0.004)
+ forest and fertiliser	0.029** (0.014)	0.020 (0.021)	0.012** (0.005)	0.011*** (0.004)
Baseline forest $\times$ year	0.027* (0.014)	0.035 (0.023)	0.009* (0.005)	0.010** (0.004)
Baseline fertiliser $\times$ year	0.025* (0.014)	0.019 (0.021)	0.013** (0.005)	0.010** (0.005)
Fixed Effects	District and State $\times$ Year			

Notes: Each cell is the parameter β of equation (5) for a 10 percentage point increase in upwind irrigation exposure. **Main** is the specification of the main text, with district and state-by-year fixed effects and no additional control, and reproduces Tables 1 and 2. Each of the two panels below has its own sample and therefore its own uncontrolled reference row, against which the rows under it should be read. Wheat and rice production are in logs and estimated on the positive observations; double cropping and irrigation intensity are bounded ratios in levels. Panel A rebuilds annual exposure from the monthly panel with the residue-burning months removed from the treatment, since dropping a month from an annual outcome would not be meaningful; the rebuilt all-months series correlates 0.998 with the published one. Panel B adds forest share of district area and fertiliser per hectare of net cropped area, contemporaneously and then as 1966–1970 baselines interacted with year. Standard errors, clustered at the district level, are in parentheses. Significance: * $p < 0.10$, ** $p < 0.05$, *** $p < 0.01$.

D.7 Spatial correlation of standard errors




Figure A27: Conley spatial standard errors over a grid of cutoffs

Notes: Figure A27 reports the estimated coefficients from specification (4) on monthly precipitation (mm) and mean daily temperature ($^{\circ}\text{C}$), upwind arm alone, with the standard error computed in five ways. Each estimated coefficient γ_m measures the *total* marginal effect of a 10 percentage point increase in upwind irrigation exposure in month m . Upwind irrigation is the mean share of irrigated area in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance, averaged over the district. Regressions are estimated at the district level and include district, month and state-by-year fixed effects, as in Figure 4. *Main* clusters by district; the four remaining series are Conley (1999) standard errors with a Bartlett taper in space, at the stated cutoff, and a three-year taper in time. The average district covers 8,900 km², an equivalent diameter of 102 km, so a cutoff below 100 km is shorter than the unit of observation and the grid begins there. The point estimate is identical across all five, so only the intervals differ. Error bars denote 90% confidence intervals.

D.8 Randomisation inference




Figure A28: Conley spatial standard errors, agricultural outcomes

Notes: Figure A28 is the agricultural counterpart of Figure A27. It reports the estimated coefficient on upwind exposure from $y_{d,t} = \beta E_{d,t}^{\uparrow} + \mu_d + \lambda_{s,t} + \varepsilon_{d,t}$ at district level, with district and state-by-year fixed effects, in the functional form of the main text: log for the four crop outcomes, estimated on the positive observations, and levels for the two bounded ratios. The point estimates are therefore exactly those of Tables 1 to 2. *Main* clusters by district, the paper's own choice; the five remaining series are Conley (1999) standard errors with a Bartlett taper in space at the stated cut-off. The point estimate is identical across all six, so only the intervals differ; estimates not significant at 10 percent are drawn at reduced opacity. Error bars denote 90% confidence intervals.




Figure A29: Randomisation inference: the treatment reassigned 500 times

Notes: Figure A29 reports the distribution of the coefficient on upwind exposure when the treatment is randomly reassigned 500 times. Each district's *entire* exposure series moves to another district, so the placebo keeps the real treatment's level, trend and serial correlation, and only the match between district and treatment is broken; every draw is then put through the paper's own specifications, equation (4) for precipitation and temperature and equation (5) for wheat and rice production, cropping intensity and irrigation intensity, in the functional form the main text reports. The grey histogram is the permutation distribution and the shaded bands its central 90 and 95 percent; the red line is the coefficient actually estimated, and the randomisation p -value is the share of permutations whose absolute coefficient is at least as large.

E Placebo




Figure A30: Directional placebo: wheat production, area, yield and irrigated area

Notes: Each panel reports the coefficient on upwind exposure beside the coefficient on downwind exposure, the two entered jointly in the same regression, so the upwind coefficient differs slightly from the tables of the main text, where it enters alone. Upwind irrigation is the average share of irrigated land in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance; the downwind arm reverses the wind. All four outcomes are in logarithm. The vertical scale is free across panels but symmetric about zero in each, so the dashed zero line sits at the same height throughout. *Main* is equation (5) on the annual panel. *5-years MA* and *10-years MA* average each district over windows of two and five years around anchor years 1970 to 2000, with district and state-by-anchor-year fixed effects. *Long-Difference* compares 1968–1972 with 1998–2002, with state fixed effects. No baseline control enters any specification. Standard errors are clustered at the district level; bars are 90 and 95 percent confidence intervals, and estimates not significant at 10 percent are drawn at reduced opacity.




Figure A31: Directional placebo: cropping and irrigation intensity

Notes: Each panel reports the coefficient on upwind exposure beside the coefficient on downwind exposure, the two entered jointly in the same regression, so the upwind coefficient differs slightly from the tables of the main text, where it enters alone. Upwind irrigation is the average share of irrigated land in pixels located 50–150 km upwind of the focal pixel, weighted inversely by distance; the downwind arm reverses the wind. *Main* is equation (5) on the annual panel. *5-years MA* and *10-years MA* average each district over windows of two and five years around anchor years 1970 to 2000, with district and state-by-anchor-year fixed effects. *Long-Difference* compares 1968–1972 with 1998–2002, with state fixed effects. No baseline control enters any specification. Standard errors are clustered at the district level; bars are 90 and 95 percent confidence intervals, and estimates not significant at 10 percent are drawn at reduced opacity. Both outcomes are bounded ratios and enter in levels.

F Counterfactual exercise

Table A17: Counterfactual exercise


	(1) Point estimate	(2) 95% interval
<i>Panel A: Canals constructed by 2010</i>		
<i>Climate</i>		
Dry-season rainfall (km ³ /year)	37	[18, 55]
mm per season	13	[6, 19]
as % of dry-season rainfall	8.9	[4.2, 13.0]
Dry-season temperature (°C)	-0.30	[-0.34, -0.26]
Monsoon rainfall (km ³ /year)	-82	[-106, -58]
<i>Agriculture</i>		
Wheat production (Mt/year)	4.9	[1.4, 7.8]
Rice production (Mt/year)	8.4	[-6.6, 19.3]
Irrigation intensity (pp)	1.71	[0.55, 2.92]
gross irrigated area (Mha)	3.2	[1.0, 5.4]
Cropping intensity (pp)	1.69	[0.68, 2.65]
<i>Value of the wheat, at current price</i>		
Gross value (bn USD/year)	1.0	
Per hectare of command area (USD/year)	19	
<i>Panel B: Canals under construction</i>		
Dry-season rainfall (km ³ /year)	48	[22, 70]
as % of dry-season rainfall	8.4	[3.9, 12.2]
Wheat production (Mt/year)	3.63	[0.94, 6.37]
<i>Value of the wheat, at current price</i>		
Gross value (bn USD/year)	0.76	
Per hectare of new command area (USD/year)	24	

Notes: National totals, evaluated over 2006–2010. Panel A sets upwind exposure to zero for the completed network, including the 5.3 million hectares with no recorded completion year. Panel B is the increment from the 32.3 million hectares under construction or completed after 2010, and reports only the two margins inside the range identifying $\hat{\beta}$. Climate uses the level-identified monthly coefficients (district-year mean exposure); agriculture uses equation (5). Intervals are the 2.5th and 97.5th percentiles of 300 district cluster bootstrap replications, the uncertainty in $\hat{\beta}$. Both panels value wheat at \$210/tonne, the mean 2007–2010 support price. Rice is not significant and is reported for completeness only.


Weather response versus adaptation. To separate the plant’s direct response to weather from what farmers did about it, I chain three stages – the effect of upwind exposure on dry-

season weather, a weather-response function for wheat yield estimated with no exposure term of its own, in the tradition of [Schlenker and Roberts \(2009\)](#), and the induced weather pushed through that function – and re-estimate the response function on Indian wheat. This agronomic channel accounts for 28–44 percent of the 3.8 percent reduced-form effect on wheat production; the remainder is consistent with the adaptation margins of Section [6.2](#).²¹

²¹[Braun and Schlenker \(2023\)](#) find a comparable split for corn and soybeans: the downwind, cooling-induced share of the yield effect is on average 27 percent of the direct on-farm irrigation effect, the same order of magnitude as the agronomic share found here.



(a) Mean temperature, Rabi



(b) Wheat area

— Observed - - - Without Irrigation . . . Ongoing Construction

Figure A32: Historical trajectories with and without upwind irrigation, extended to 2035: temperature and wheat area

Notes: As the national historical-trajectory figure of Section 8.




Figure A33: What the projects that came after 2010 would deliver, state by state
Notes: Points are the increment each state would receive; bars are 95 percent district cluster bootstrap intervals over 300 draws. **All India** in red.




Figure A34: What the standing canal network delivers to each state today
Notes: As Figure A33, for the network already standing, averaged over 2006–2010.




Figure A35: India with and without upwind irrigation, state by state

Notes: As the national historical-trajectory figure of Section 8, for the nine states with the largest effect. About 0 percent of Punjab's district-months and 7 percent of Haryana's have a negative counterfactual.